

Version 1.03

UCSD CSE 30


Computer Organization and Systems Programming

Numbers Data and Memory

Lecture 4 – Oct 4, 2022

Keith Muller

Cray XMP-1982



Review: Decimal Numbering

- Decimal is base 10
 - from “decem” (Latin) \Rightarrow Ten Characteristics
- **Ten** symbols (Why?)
0 1 2 3 4 5 6 7 8 9
- **How do we represent larger numbers?**
 - **Large numbers are a sequence of digits**
 - Each digit is one of the available symbols
 - n-digit numbers as coefficients in a polynomial



$$d_{n-1} d_{n-2} \dots d_1 d_0 = d_{n-1} \cdot 10^{n-1} + d_{n-2} \cdot 10^{n-2} + \dots + d_1 \cdot 10 + d_0$$

Diagram illustrating the polynomial representation of a decimal number. The equation shows the sequence of digits $d_{n-1} d_{n-2} \dots d_1 d_0$ equal to the sum of each digit multiplied by its corresponding power of 10. A box labeled "Digit position" points to the exponents $n-1, n-2, \dots, 1, 0$. A box labeled "symbol value" points to the digits $d_{n-1}, d_{n-2}, \dots, d_1, d_0$.

- **Examples**
 - 7061 in decimal (base 10)
 - $7061_{10} = (7 \times 10^3) + (0 \times 10^2) + (6 \times 10^1) + (1 \times 10^0)$

Review: Binary Numbering

- Binary is base 2
 - *adjective*: being in a **state of one of two mutually exclusive conditions** such as **on** or **off**, **true** or **false**, **molten** or **frozen**, **presence** or **absence** of a signal
 - From Late Latin *bīnārius* (“consisting of two”)
- **Two** symbols:
0 1
- Numbers in C starting with **0b** are binary
- Example: What is **0b**110 in base 10?
 - **0b**110 = $110_2 = (1 \times 2^2) + (1 \times 2^1) + (0 \times 2^0) = 6_{10}$
- A **bit** is a **single binary digit**
- A **byte** is an **8-bit** value

powers of two



$$\text{Unsigned binary Number} = \sum_{i=0}^{n-1} b_i \times 2^i = b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

Review: Octal Numbering

- **Eight** symbols

0, 1, 2, 3, 4, 5, 6, 7

Notice that we no longer use 8 or 9

- Base comparison:

- **Base 10**: 0, 1, 2, 3, 4, 5, 6, 7, **8**, 9, 10, 11, 12..

- **Base 8**: 0, 1, 2, 3, 4, 5, 6, 7, **10**, 11, 12, 13, 14..

- Numbers in C starting with a **0**: 07061, are octal

- Example: What is 07061₈ in base 10?

- **07061**₈ = (7 × 8³) + (0 × 8²) + (6 × 8¹) + (1 × 8⁰) = 3633₁₀

subscript indicates base

powers of eight

in C leading 0 indicates octal

$$\text{Unsigned octal Number} = \sum_{i=0}^{n-1} b_i \times 8^i = b_{n-1}8^{N-1} + b_{n-2}8^{N-2} + \dots + b_18^1 + b_08^0$$



Review: Hexadecimal Numbering

- hexadecimal is base 16
 - From “hexa” (Ancient Greek ἑξά-) \Rightarrow six
 - and from “decem” (Latin) \Rightarrow ten

- **Sixteen** symbols

0 1 2 3 4 5 6 7 8 9 a b c d e f

- Numbers in C starting with **0x** are hexadecimal

- $16_{10} = \text{0x}10_{16}$

- Example: What is **0xa5** in base 10?

- $\text{0xa5} = \text{a5}_{16} = (10 \times 16^1) + (5 \times 16^0) = 165_{10}$

- **Hexadecimal** numbers are **very commonly used** in programming to express binary values

- Imagine the difficulty in correctly expressing a 64-bit binary value in your code

$$\text{Unsigned Hex Number} = \sum_{i=0}^{n-1} b_i \times 16^i = b_{n-1}16^{n-1} + b_{n-2}16^{n-2} + \dots + b_116^1 + b_016^0$$



Number Base Overview (as written in C)

- Decimal is base 10, Hexadecimal is base 16, and octal is base 8
- **Octal digits** have 8 values 0 – 7 (written in C as **00** – **07**, careful **073** is octal = 59 in decimal)
- **Hex digits** have 16 values 0 - 9 a - f (written in C as **0x0** – **0xf**)
- No standard prefix in C for binary (most use **hex**) – gcc (compiler) allows **0b** prefix **others might not**

Hex digit Octal digit	0x0 00	0x1 01	0x2 02	0x3 03	0x4 04	0x5 05	0x6 06	0x7 07
Decimal value	0	1	2	3	4	5	6	7
Binary value	0b0000	0b0001	0b0010	0b0011	0b0100	0b0101	0b0110	0b0111
Hex digit Octal digit	0x8 010	0x9 011	0xa 012	0xb 013	0xc 014	0xd 015	0xe 016	0xf 017
Decimal value	8	9	10	11	12	13	14	15
Binary value	0b1000	0b1001	0b1010	0b1011	0b1100	0b1101	0b1110	0b1111

Binary <---> Hexadecimal Equivalences

- Hex → Binary: $16^1 = 2^4$ 1 digit hex = 4 digits binary
 1. Replace hex digits with binary digits
 2. drop **leading zeros**
 - Example: 0x2d to binary
 - 0x2 is 0b0010, 0xd is 0b1101
 - Drop two leading zeros, answer is 0b101101
- Binary → Hex: $2^4 = 16^1$
 1. Pad with enough **leading zeros** until number of digits is a multiple of 4
 2. replace each **group of 4** with the **HEX equivalent**
 - Example: 0b101101
 - **Pad on the left** to: 0b 0010 1101
 - Replace to get: 0x2d

Base 10	Base 2	Base 16
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
10	1010	a
11	1011	b
12	1100	c
13	1101	d
14	1110	e
15	1111	f

Hex to Binary (group 4 bits per digit from the right)

- Each Hex digit is 4 bits in base 2 $16^1 = 2^4$

0x f a 5 3

1111 1010 0101 0011

0b1111101001010011

binary start with a 0b in C

Binary to Hex (group 4 bits per digit from the right)

- 4 binary bits is one Hex digit $2^4 = 16^1$

0b 0110 1010 0011 1111

6 a 3 f

0x6a3f

hex start with 0x in C



Binary to Octal (group 3 bits per digit from the right)

- 3 binary bits is one Octal digit $2^3 = 8^1$


0b 0 110 101 000 111 111
0 6 5 0 7 7

065077

octal start with a 0 in C

Octal to Binary (group 3 bits per digit from the right)

- One Octal digit is three binary digits $2^3 = 8^1$

0	1	7	5	1	2	3
						
1	1	1	1	0	1	0
0	0	0	1	0	1	0
0	1	0	0	1	1	1



0b1111101001010011

 binary start with a 0b in C

Looking at the Powers of Two

$$\text{Unsigned binary Number} = b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

Base 10	$b_5 \cdot 32$	$b_4 \cdot 16$	$b_3 \cdot 8$	$b_2 \cdot 4$	$b_1 \cdot 2$	b_0
5	0	0	0	1	0	1
10	0	0	1	0	1	0
20	0	1	0	1	0	0
40	1	0	1	0	0	0

- Multiple – left shift 
- Divide – right shift 

- Divide by 10 (whole numbers) – shift left by a digit
 - $123/10 = 12$
- Multiply by 10 (whole numbers) – shift right by a digit
 - $123 \times 10 = 1230$

Multiply By Base: Shift Left 1 digits = Multiply by 2




Most significant
Digit

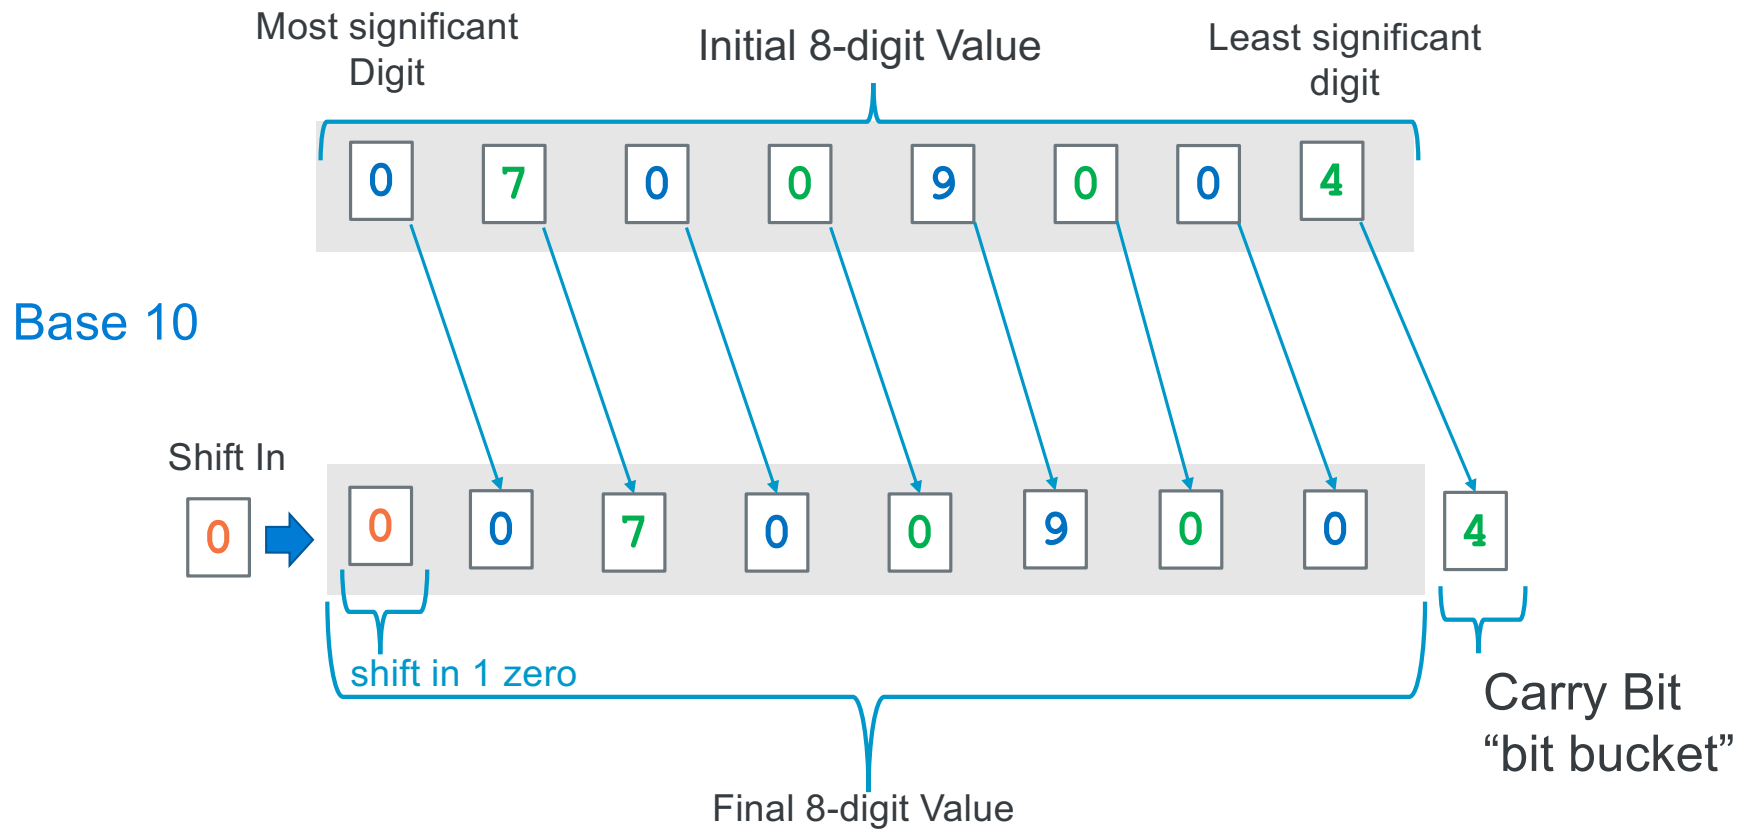

Initial 8-digit Value

Least significant
Digit

Base 10



Divide by Base: Shift Right 1 Digit = Divide by 2




Use a Right Bit Shift: Unsigned Decimal To Unsigned Binary

Here is an Algorithmic approach to convert unsigned numbers to binary

Perform a sequence of **divisions** (**right shift**) to isolate the remainder

$$\text{Unsigned binary Number} = b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

$$\text{Unsigned Binary Number} = 2 \times (\dots (2 \times (2 \times b_{n-1} + b_{n-2})) + \dots + b_1) + b_0$$



Unsigned Decimal to Unsigned Binary Conversion

	dividend 249	Quotient	Remainder	Bit Position
➡	249/2	124	➡ 1	b0
➡	124/2	62	➡ 0	b1
➡	62/2	31	➡ 0	b2
➡	31/2	15	➡ 1	b3
➡	15/2	7	➡ 1	b4
➡	7/2	3	➡ 1	b5
➡	3/2	1	➡ 1	b6
➡	1/2	0	➡ 1	b7


$$249(\text{base } 10) = b_7 b_6 b_5 b_4 b_3 b_2 b_1 b_0 = 0b1111001$$

$$11111001 = (1 \times 128) + (1 \times 64) + (1 \times 32) + (1 \times 16) + (1 \times 8) + 1 = 249$$

Left Bit Shift & add: Unsigned Binary to Unsigned Decimal

$$b_{n-1}2^{N-1} + b_{n-2}2^{N-2} + \dots + b_12^1 + b_02^0$$

- Base conversion via a sequence of n multiplications (left shift) and n additions
 - 111 base 2 $\rightarrow ((1 \times 2 + 1) \times 2) + 1$
- Alternatively, you can memorize and use the positional weights to convert



$$11101 \text{ base 2} = (1 \times 16) + (1 \times 8) + (1 \times 4) + (0 \times 2) + (1 \times 1) = 29$$

Unsigned Binary to Unsigned Decimal Conversion

What is $0 \overset{b_7}{1} \overset{b_6}{1} \overset{b_5}{0} \overset{b_4}{0} \overset{b_3}{1} \overset{b_2}{0} \overset{b_1}{1} \overset{b_0}{1}_{(\text{base } 2)}$ in decimal (N)?

	Product Shift Left	Addend	Bit Position	Product
→	0	+ 0	b7	0
→	0	+ 1	b6	1
→	2 x 1 = 2	+ 1	b5	3
→	2 x 3 = 6	+ 0	b4	6
→	2 x 6 = 12	+ 0	b3	12
→	2 x 12 = 24	+ 1	b2	25
→	2 x 25 = 50	+ 0	b1	50
→	2 x 50 = 100	+ 1	b0	101

$101_{(\text{base } 10)} = (1 \times 64) + (1 \times 32) + (1 \times 4) + 1$ (checking the conversion)

Version 1.03

UCSD CSE 30


Computer Organization and Systems Programming

Numbers Data and Memory

Lecture 4 – Oct 4, 2022

Keith Muller

Cray XMP-1982



Different Type of Numbers each have a Fixed # of Bits


Spanning one or more contiguous bytes of memory

C Data Type	AArch-32 contiguous Bytes
char (arm unsigned)	1
short int	2
unsigned short int	2
int	4
unsigned int	4
long int	4
long long int	8
float	4
double	8
long double	8
pointer *	4

Byte 8-bit integer uses 1 byte




Half Word 16-bit integer uses 2 bytes



most significant bit (largest power of 2)

least significant byte

Word 32-bit integer uses 4 bytes




most significant byte

most significant bit (smallest power of 2)

Unsigned Integers (positive numbers) Impact of Fixed # of Bits

- 4 bits is $2^4 = \text{ONLY } 16$ distinct values
- **Modular** (C operator: `%`) or **clock math**
 - Numbers start at 0 and “wrap around” after 15 and go back to 0
- Keep **adding** 1
 - wraps (**clockwise**)
 - 0000 \rightarrow 0001 ... \rightarrow 1111 \rightarrow 0000
- Keep **subtracting** 1
 - wraps (**counter-clockwise**)
 - 1111 \rightarrow 1110 ... \rightarrow 0000 \rightarrow 1111
- Addition and subtraction use normal “**carry**” and “**borrow**” rules, just operate in binary



Addition

- Add two numbers in decimal

$$\begin{array}{r} 27 \\ + 20 \\ \hline 47 \end{array}$$

+1 carry = 2

↓

$$\begin{array}{r} 12 \\ + 78 \\ \hline 90 \end{array}$$

- Adding Single digits in binary

- $0 + 0 = 0$
- $0 + 1 = 1$
- $1 + 0 = 1$
- $1 + 1 = 0$ (with 1 carry) or $10_{\text{base } 2}$

Unsigned Binary Number: Addition in 8 bits

Be Aware in Binary

$1 + 1 = 10$ base 10: $(1 + 1 = 2)$

$1 + 1 + 1 = 11$ base10: $(1 + 1 + 1 = 3)$

Carry
Bit

carries

0 0 1 0 0 0 1 1

+

1 0 1 0 0 0 0 1

161

0 0 1 1 0 0 1 1

51

sum

1 1 0 1 0 1 0 0

212

Subtraction

- Add two numbers in decimal

		<div>-1 borrow = 5</div>	<div>+10 borrow = 12</div>
		↓	↓
	2 7	6 2	
-	<u>1 0</u>	+	<u>2 3</u>
	1 7		3 9

- Subtracting Single digits in binary
 - $0 - 0 = 0$
 - $0 - 1 = 1$ (with a +2 (a 1) from the column to the left)
 - $1 - 0 = 1$
 - $1 - 1 = 0$

Unsigned Binary Number: Subtraction in 8 bits

borrows

$$\begin{array}{r} 10100001 \\ - 00110011 \\ \hline \end{array}$$

difference

Be Aware in Binary

$$1 - 1 = 0$$

$$10 - 1 = 1 \text{ base 10: } (2 - 1 = 1)$$

Unsigned Binary Number: Subtraction in 8 bits

Build of previous slide – note warning

borrows		10	10	10	10	10	10			
		0	1	0	1	1	1	0	1	161
	–	0	0	1	1	0	0	1	1	51
		<hr/>								
difference		0	1	1	0	1	1	1	0	110

Be Aware in Binary


$$1 - 1 = 0$$

$$10 - 1 = 1 \text{ base 10: } (2 - 1 = 1)$$

NOTICE

This slide is performing the Subtraction shown in the previous slide using PowerPoint builds when this slide is viewed in a pdf it will look incorrect

Overflow: Going Past the Boundary Between max and min



Overflow: Occurs when an arithmetic result (from addition or subtraction for example) is **is more than min** or **max** limits

C (and Java) ignore overflow exceptions

- You end up with a bad value in your program and absolutely no warning or indication... **happy debugging!....**

Unsigned Integer Number Overflow: Addition in 8 bits

Carry Bit

carries

only 8 bits for numbers in this example carry bit is always dropped from result

+

sum

1	1	1	1	1	1	1	1	
1	0	1	0	0	0	0	1	161
0	1	0	1	1	1	1	1	95
<hr/>								
0	0	0	0	0	0	0	0	256

Rule: When Carry Bit $\neq 0$, overflow has occurred for unsigned integers!

Overflow: Unsigned Values 4-bit limit

Addition Overflow:

$$\begin{array}{r} 15 \\ + 2 \\ \hline 17 \end{array}$$

only 4 bits for numbers in this example

carry bit is dropped from the result


$$\begin{array}{r} 1111 \\ + 0010 \\ \hline 10001 \\ \text{oops } 1 \end{array}$$

Subtraction Overflow:

$$\begin{array}{r} 1 \\ - 2 \\ \hline -1 \end{array}$$

$$\begin{array}{r} 0001 \\ - 0010 \text{ (counter clockwise)} \\ \hline 1111 \\ \text{oops } 15 \end{array}$$

Overflow: Occurs when an arithmetic result is **exceeds** the min or max limits



Unsigned Binary Multiplication example

$$\begin{array}{r}
 a = 01010011 \quad (83) \\
 \times b = 0000010 \quad (2) \\
 \hline
 00000000 \\
 01010011 \\
 \hline
 010100110 \quad (166)
 \end{array}$$

Since we are multiplying by a power of two we can shift left by 2 (zero insert at LSB)


$$\begin{array}{r}
 a = 01010011 \quad (83) \\
 \swarrow \swarrow \swarrow \swarrow \swarrow \swarrow \swarrow \swarrow \\
 010100110 \quad (166)
 \end{array}$$

$$\begin{array}{r}
 a = 01010011 \quad (83) \\
 \times b = 0000101 \quad (5) \\
 \hline
 01010011 \\
 00000000 \\
 01010011 \\
 \hline
 011001111 \quad (415!!)
 \end{array}$$


Exceeds what we can represent in 8 bits (255)

Unsigned Binary Number Divide

Consider $12/3 = 4$



Consider $12/2 = 6$




Since we are dividing by a power of two we can shift right by 1 (zero insert at MSB)



Problem: How to Encode Both Positive and Negative Integers

- How do we represent the negative numbers within a fixed number of bits?
 - Allocate some bit patterns to negative and others to positive numbers (and zero)
- 2^n distinct bit patterns to encode positive and negative values
- Unsigned values:** $0 \dots 2^n - 1$ ← -1 comes from counting 0 as a "positive" number
- Signed values:** $-2^{n-1} \dots 2^{n-1} - 1$ (dividing the range in ~ half including 0)
- On a number line (below):** 8-bit integers – signed and unsigned (e.g., `char` in C)




Same "width" (same number of encodings), just shifted in value

Negative Integer Numbers: Sign + Magnitude Method



- Use the **M**ost **S**ignificant **B**it as a sign bit
 - 0 as the MSB represents positive numbers
 - 1 as the MSB represents negative numbers
- **Two** (oops) representations for **zero**: 0000, 1000
- Tricky Math (must handle sign bit independently)
 - **Positive and Negatives “increment” (+1) in the opposite directions!**



Signed Magnitude Examples (Sign bit is always MSB)

0 110
positive 6

1 011
negative 3

Examples (4 bits):


1 000 = -0	0 000 = 0
1 001 = -1	0 001 = 1
1 010 = -2	0 010 = 2
1 011 = -3	0 011 = 3
1 100 = -4	0 100 = 4
1 101 = -5	0 101 = 5
1 110 = -6	0 110 = 6
1 111 = -7	0 111 = 7

0 00000000
positive 0

1 0001100
negative 12

2's Complement Signed Integer Method


- Positive numbers encoded same as unsigned numbers
- All **negative values** have a **one in the leftmost bit**
- All **positive values** have a **zero in the leftmost bit**
 - This implies that 0 is a positive value
- **Only one zero**
- **For n bits, Number range is $-(2^{n-1})$ to $+(2^{n-1} - 1)$**
 - Negative values “go 1 further” than the positive values
- Example: the range for 8 bits:
-128, -127, .. 0, .. 126, +127
- Example the range for 32 bits:
-2147483648 .. 0, .. +2147483647
- *Arithmetic is the same as with unsigned binary!*



Two's Complement: The MSB Has a *Negative Weight*

$$2's\ Comp = -b_{n-1}2^{n-1} + b_{n-2}2^{n-2} + \dots + b_12^1 + b_02^0$$

b_{n-1} weight is (-2^{n-1}) , all other bits: have positive weights $(+2^i)$



- 4-bit ($w = 4$) weight = $-2^{4-1} = -2^3 = -8$
 - 1010_2 **unsigned**:
 $1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 = 10$
 - 1010_2 **two's complement**:
 $-1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 = -8 + 2 = -6$
 - -8 in **two's complement**:
 $1000_2 = -2^3 + 0 = -8$
 - -1 in **two's complement**:
 $1111_2 = -2^3 + (2^3 - 1) = -8 + 7 = -1$

Negation Of a Two's Complement Number (Method 1)

$$\begin{array}{r}
 7 = 0111 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 1000 \\
 \text{add } 1 \quad + \quad \underline{\quad 1 \quad} \\
 -7 \quad \quad 1001
 \end{array}$$

$$\begin{array}{r}
 -7 = 1001 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 0110 \\
 \text{add } 1 \quad + \quad \underline{\quad 1 \quad} \\
 7 \quad \quad 0111
 \end{array}$$

$$-x == \sim x + 1;$$

$$\begin{array}{r}
 7 = \quad \quad 0111 \\
 -7 = \quad + \quad \underline{1001} \\
 \text{(discard carry)} \quad 0000
 \end{array}$$




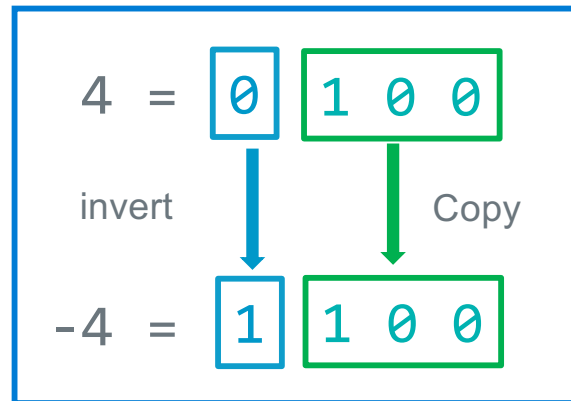
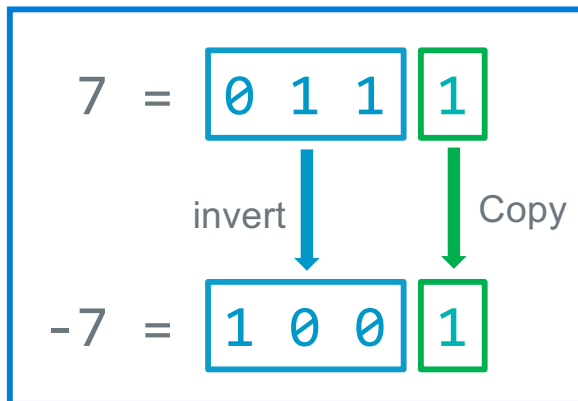
$$\begin{array}{r}
 1 = 0001 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 1110 \\
 \text{add } 1 \quad + \quad \underline{\quad 1 \quad} \\
 -1 \quad \quad 1111
 \end{array}$$

$$\begin{array}{r}
 -1 = 1111 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = 0000 \\
 \text{add } 1 \quad + \quad \underline{\quad 1 \quad} \\
 1 \quad \quad 0001
 \end{array}$$

$$\begin{array}{r}
 -8 = 1000 \\
 \downarrow \downarrow \downarrow \downarrow \\
 \text{invert} = \underline{0111} \\
 \text{add } 1 \quad + \quad \underline{\quad 1 \quad} \\
 -8 \quad \quad 1000 \text{ oops!}
 \end{array}$$

Negation of a Two's Complement Number (Method 2)

1. **copy unchanged** right most bit containing a 1 and all the 0's to its right
2. Invert all the bits to the left of the right-most 1



Signed Decimal to Two's Complement Conversion

	dividend -102	Quotient	Remainder	Bit Position
➡	102/2	51	➡ 0	b0
➡	51/2	25	➡ 1	b1
➡	25/2	12	➡ 1	b2
➡	12/2	6	➡ 0	b3
➡	6/2	3	➡ 0	b4
➡	3/2	1	➡ 1	b5
➡	1/2	0	➡ 1	b6
➡	0/2	0	➡ 0	b7

102(base 10) = $b_7 b_6 b_5 b_4 b_3 b_2 b_1 b_0$ = 0b0110 0110

Get the two complement of 01100110 is 10011010



Two's Complement to Signed Decimal Conversion - Positive

What is $b_7 \ b_6 \ b_5 \ b_4 \ b_3 \ b_2 \ b_1 \ b_0$
 0 1 1 0 0 1 0 1_(base 2) in decimal (N)?

Signed Bit Bias	Bit	Bit Position	Bias
$-2^{W-1} = -2^{8-1} = -128$	x 0	b7	0 ←
Product Shift Left	Addend	Bit Position	Product
0	+ 1	b6	1
2 x 1 = 2	+ 1	b5	3
2 x 3 = 6	+ 0	b4	6
2 x 6 = 12	+ 0	b3	12
2 x 12 = 24	+ 1	b2	25
2 x 25 = 50	+ 0	b1	50
2 x 50 = 100	+ 1	b0	SUM = 101
		Bias + SUM:	0 + 101 = 101

Two's Complement to Signed Decimal Conversion - Negative

What is $\overset{b_7}{1} \overset{b_6}{1} \overset{b_5}{1} \overset{b_4}{0} \overset{b_3}{0} \overset{b_2}{1} \overset{b_1}{0} \overset{b_0}{1}_{(\text{base } 2)}$ in decimal (N)?

Signed Bit Bias	Bit	Bit Position	Bias
$-2^{W-1} = -2^{8-1} = -128$	x 1	b7	-128
Product Shift Left	Addend	Bit Position	Product
0	+ 1	b6	1
2 x 1 = 2	+ 1	b5	3
2 x 3 = 6	+ 0	b4	6
2 x 6 = 12	+ 0	b3	12
2 x 12 = 24	+ 1	b2	25
2 x 25 = 50	+ 0	b1	50
2 x 50 = 100	+ 1	b0	SUM = 101
		Bias + SUM:	-128 + 101 = -27

Two's Complement Addition and Subtraction

- **Addition:** just add the two number directly
- **Subtraction:** you can convert to addition: $\text{difference} = \text{minuend} - \text{subtrahend}$
 $\text{difference} = \text{minuend} + 2\text{'s complement}(\text{subtrahend})$

	Count	0	0	0	0	0	0	1	1
x	=	0	1	0	1	0	0	1	1
y	=	0	0	0	0	1	0	1	1
x + y	=	0	1	0	1	1	1	1	0

$$\begin{array}{r} \mathbf{x} = 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \\ \mathbf{y} = 0 \ 0 \ 0 \ 0 \ 1 \ 0 \ 1 \ 1 \\ \hline \mathbf{x-y} = 0 \ 1 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \end{array}$$



2's complement first and then add


$$\begin{array}{r} \mathbf{x} = 0 \ 1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 1 \\ + \ (-\mathbf{y}) = 1 \ 1 \ 1 \ 1 \ 0 \ 1 \ 0 \ 1 \\ \hline \mathbf{x} - \mathbf{y} = \mathbf{x} + (-\mathbf{y}) = 0 \ 1 \ 0 \ 0 \ 1 \ 0 \ 0 \ 0 \end{array}$$

Two's Complement Positive Overflow

- **4-bit** Two's complement numbers (positive overflow)

Cout	0	1	0	0	
	0	1	0	1	5
+	0	1	1	0	6
<hr/>					
	1	0	1	1	-5
					!= 11

signed numbers: overflow occurs if
operands have same sign and result's sign is different




Overflow: Occurs when an arithmetic result is beyond the min or max limits

Two's Complement Negative Overflow

- **4-bit** Two's complement numbers (negative overflow)

	Cout	1	0	1	1	
		1	0	0	1	-7
		+	1	0	1	-5
carry bit is dropped from result		0	1	0	0	+4
						!= -12



signed numbers: overflow occurs if
operands have same sign and result's sign is different

Overflow: Occurs when an arithmetic result is beyond the min or max limits

Summary: When Does Overflow Occur

Operand 1
+ Operand 2
Result

Operand 1 Sign	Operand 2 Sign	Is overflow Possible?
+	+	YES
-	-	YES
+	-	NO
-	+	NO

Sign Extension in C: Type casts

- Convert from smaller to larger integral data types
- C and Java automatically performs sign extension
- Example (on pi-cluster with 32-bit int and 16-bit short)

```
#include <stdlib.h>
#include <stdio.h>
int main(void)
{
    signed char c = -1;
    signed int i = c;
    unsigned char d = 1;
    unsigned int j = d;
    printf("c decimal = %hd\n", c);
    printf("c = 0x%hhx\n", c);
    printf("i decimal = %d\n", i);
    printf("i = 0x%x\n", i);
    printf("\nd decimal = %hd\n", d);
    printf("d = 0x%hhx\n", d);
    printf("j decimal = %d\n", j);
    printf("j = 0x%x\n", j);
    return EXIT_SUCCESS;
}
```

```
./a.out
c decimal = -1
c = 0xff
i decimal = -1
i = 0xffffffff

d decimal = 1
d = 0x1
j decimal = 1
j = 0x1
```

Sign Extension (how type promotion works)

- Sometimes you need to work with integers encoded with different number of bits

8 bits (char) → (16 bits) short → (32 bits) int


- Sign extension increases the number of bits:** n -bit wide signed integer X , **EXPANDS** to a **wider** n -bit + k -bit signed integer X' where **both have the same value**

Unsigned

- Just add leading zeroes to the left side



Two's Complement Signed:

- If **positive**, add leading **zeroes on the left**
 - Observe: Positive stay positive
- If **negative**, add **leading ones on the left**
 - Observe: Negative stays negative



Example: Two's Complement Sign or bit Extension - 1

- Adding 0's in front of a positive number does not change its value




Example: Two's Complement Sign or bit Extension -2

- Adding 1's if front of a negative number does not change its value

```

    7 = 0111
        ↓ ↓ ↓ ↓
invert = 1000
add 1  +   1
      -----
    -7  1001

```

$-7 = 1001$
 extend to
 8 bits

 11111001

$$\begin{aligned} 1001 &= -8 + 1 = -7 \\ \text{1111}1001 &= \\ (-128 + 64 + 32 + 16 + 8) + 1 &= \\ = -8 + 1 &= -7 \end{aligned}$$


7 = 00000111
invert = 11111000
add 1 + 1
-7 = 11111001

Sign Extension Under Different representations

How to sign extend
this bit pattern?


0xd6

unsigned




0x00d6

signed
magnitude



0x8056

two's
complement



0xffd6

Summary: Min, Max Values: Unsigned and Two's Complement

Two's Complement → Unsigned for n bits

- **Unsigned Value Range**

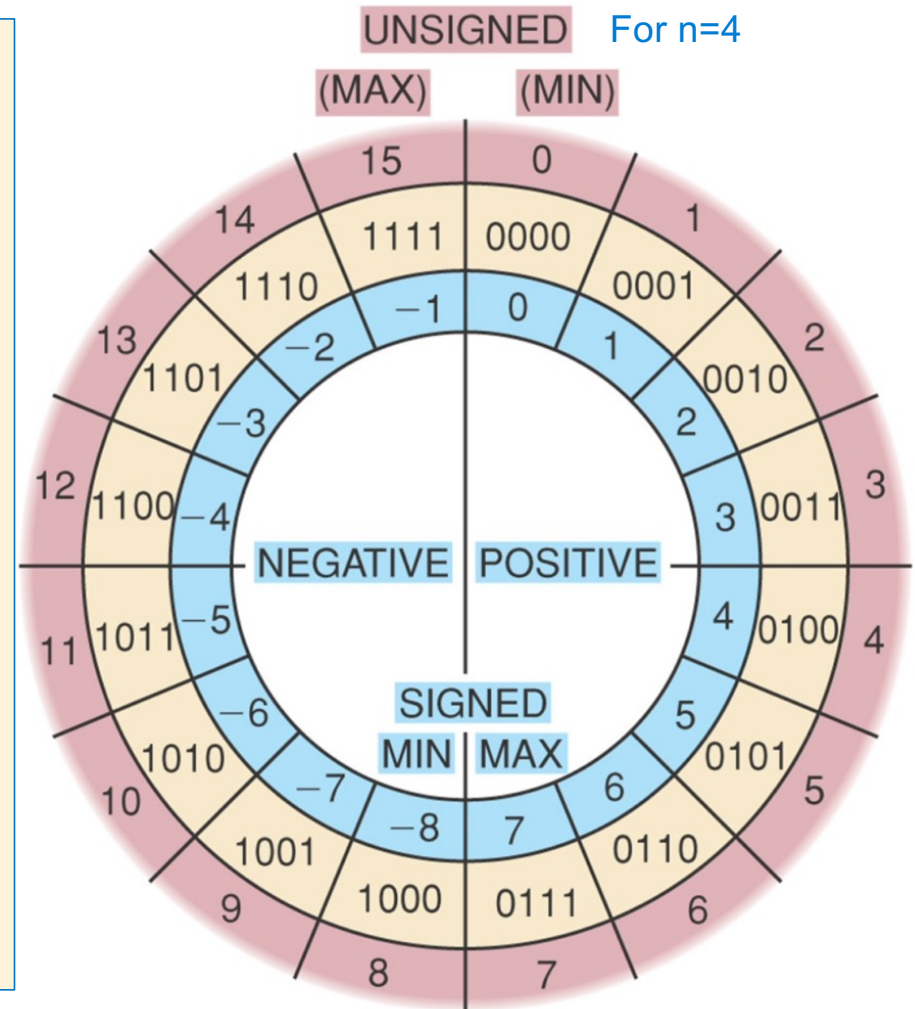
$$\begin{aligned} \text{UMin} &= 0b00\dots00 \\ &= 0 \end{aligned}$$

$$\begin{aligned} \text{UMax} &= 0b11\dots11 \\ &= 2^n - 1 \end{aligned}$$

- **Two's Complement Range**

$$\begin{aligned} \text{SMin} &= 0b10\dots00 \\ &= -2^{n-1} \end{aligned}$$

$$\begin{aligned} \text{SMax} &= 0b01\dots11 \\ &= 2^{n-1} - 1 \end{aligned}$$



Variables: Size

- **Integer types**

- `char`, `int`

- **Floating Point**

- `float`, `double`

- **Modifiers for each base type**

- `short` [int]
- `long` [int, double]
- `signed` [char, int]
- `unsigned` [char, int]
- `const`: variable read only

- **char type**

- One byte in a byte addressable memory
- **Signed** vs **Unsigned** Char implementations
- **Be careful** `char` is unsigned on arm and signed on other HW like intel

C Data Type	AArch-32 contiguous Bytes	AArch-64 contiguous Bytes	printf specification
<code>char</code> (arm unsigned)	1	1	%c
<code>short int</code>	2	2	%hd
<code>unsigned short int</code>	2	2	%hu
<code>int</code>	4	4	%d / %i
<code>unsigned int</code>	4	4	%u
<code>long int</code>	4	8	%ld
<code>long long int</code>	8	8	%lld
<code>float</code>	4	4	%f
<code>double</code>	8	8	%lf
<code>long double</code>	8	16	%Lf
<code>pointer *</code>	4	8	%p

size of a pointer is the word size

Fixed size types in C

- Sometimes programs need to be written for a particular range of integers or for a particular size of storage, regardless of what machine the program runs on
- We will need to do this in PA6 and PA7
- In the file `<stdint.h>` the following fixed size types are defined for use in these situations:

Signed Data types	Unsigned Data types	Exact Size
<code>int8_t</code>	<code>uint8_t</code>	8 bits (1 byte)
<code>int16_t</code>	<code>uint16_t</code>	16 bits (2 bytes)
<code>int32_t</code>	<code>uint32_t</code>	32 bits (4 bytes)
<code>int64_t</code>	<code>uint64_t</code>	64 bits (8 bytes)


Example: Limits of Some Types on 32-bit ARM

type	Smallest	largest
unsigned char	0	255
char (ARM - unsigned)	0	255
char (INTEL - signed)	-128	127
unsigned short	0	65,535
short	-32,768	32,767
uint32_t	0	$2^{32} - 1$
int32_t	-2^{31}	$2^{31} - 1$


- **WARNING:** If you want to force a char to be signed or unsigned use either
 - signed char or unsigned char

Fractional Binary Numbers (Software not Hardware)

Binary	Decimal
2^{-1}	0.5
2^{-2}	0.25
2^{-3}	0.125
2^{-4}	0.0625



- “**Binary Point,**” like **decimal point**, signifies boundary between **integer** and **fractional parts**
- Bits to right of “binary point” represent **fractional powers of 2**
- Example: $10.1010_2 = 1 \times 2^1 + 1 \times 2^{-1} + 1 \times 2^{-3} = 2.625_{10}$



Fixed Point Binary Number Divide

Consider $10/4 = 2.5$

$$\begin{array}{r}
 \begin{array}{ccccccc}
 & 0 & 0 & 1 & 0 & . & 1 \\
 100 & \overline{) 1} & 0 & 1 & 0 & . & 0 \\
 & -1 & 0 & 0 & & & \\
 & \hline
 & & 0 & 1 & 0 & & \\
 & & -0 & 0 & 0 & & \\
 & & \hline
 & & & 1 & 0 & 0 & \\
 & & & -1 & 0 & 0 & \\
 & & & \hline
 & & & & & 0 &
 \end{array}
 \end{array}$$

Consider $10/8 = 1.25$

$$\begin{array}{r}
 \begin{array}{ccccccc}
 & 0 & 0 & 0 & 1 & . & 0 & 1 \\
 1000 & \overline{) 1} & 0 & 1 & 0 & . & 0 & 0 \\
 & -1 & 0 & 0 & 0 & & & \\
 & \hline
 & & 0 & 1 & 0 & 0 & & \\
 & & -0 & 0 & 0 & 0 & & \\
 & & \hline
 & & & 1 & 0 & 0 & 0 & \\
 & & & -1 & 0 & 0 & 0 & \\
 & & & \hline
 & & & & & & 0 &
 \end{array}
 \end{array}$$

Fixed Point Binary Number Limitations

- Limitation #1


- Can only exactly represent numbers of the form $x/2^k$
 - Other rational numbers have repeating bit representations

Value	Representation
1/3	0.0101010101 [01] ... ₂
1/5	0.001100110011 [0011] ... ₂
1/10	0.0001100110011 [0011] ... ₂

- Limitation #2

- Just one setting of binary point within the w bits
- Limited range of numbers (very small values? very large?)

Scientific Notation Binary



- Computer hardware that supports this is called **floating point hardware** due to the “floating” of the binary point
- Declare such variable in C as `float` (or `double`)

Floating Point Representation

- Analogous to scientific notation
- In Decimal:
 - Not 12000000, but 1.2×10^7 In C: 1.2e7
 - Not 0.0000012, but 1.2×10^{-6} In C: 1.2e-6
- In Binary:
 - Not 11000.000, but 1.1×2^4
 - Not 0.000101, but 1.01×2^{-4}

Normalized Scientific Notation

- Convert from **scientific notation** to fixed **binary point**
- Perform the multiplication by shifting the decimal until the exponent disappears

Binary	Decimal
2^{-1}	0.5
2^{-2}	0.25
2^{-3}	0.125
2^{-4}	0.0625

- Example: $1.011_2 \times 2^4 = 10110_2 = 22_{10}$
- Example: $1.011_2 \times 2^{-2} = 0.01011_2 = 0.34375_{10}$
- Convert from **binary point** to **normalized scientific notation**
 - Distribute out exponents until binary point is to the right of a single digit
 - Example: $1101.001_2 = 1.101001_2 \times 2^3$

Encoding Fractions Observations

In Base 2:

$$10.1 \times 2^5 = 1.01 \times 2^6$$

$$1011.1 \times 2^5 = 1.0111 \times 2^8$$

$$0.110 \times 2^5 = 1.10 \times 2^4$$

Normalizing with base 2 :

adjust so there *always* a 1 to the **left of the decimal point!**

this 1 is **called the hidden bit** as we do not have use a bit to store it since it is there in every normalized mantissa


- Adjust x to always be in the format **1.XXXXXXXXXX...** (**fraction is normalized**)
- Fraction portion ONLY **encodes** what is *to the right* of the decimal point
- “Hidden bit” allows number to have **One additional digit for increased precision**

Fraction encoding is **1.[FRACTION BINARY DIGITS]**

Floating Point Numbers: Implementation Approach


- Supports a wide range of numbers
- Flexible “floating” decimal point
- Represent scientific notation numbers like 1.202×10^6

$$(-1)^S M 2^E$$



- **Sign bit** (a single bit): 0 positive, 1 negative
- **Exponent:** encoding of E above (it is NOT E directly represented in binary)
- **Fraction:** encoding of M above (it is NOT M directly represented in binary)

Floating Point Number in a Byte (Not A Real Format)



- **Mantissa encoding:** = 1.[xxxx] encoded as an unsigned value
 - 4 bits = 16 values + leading digit is always a 1
- **Exponent encoding:** 3 bits encoded to represent both – and + exponents

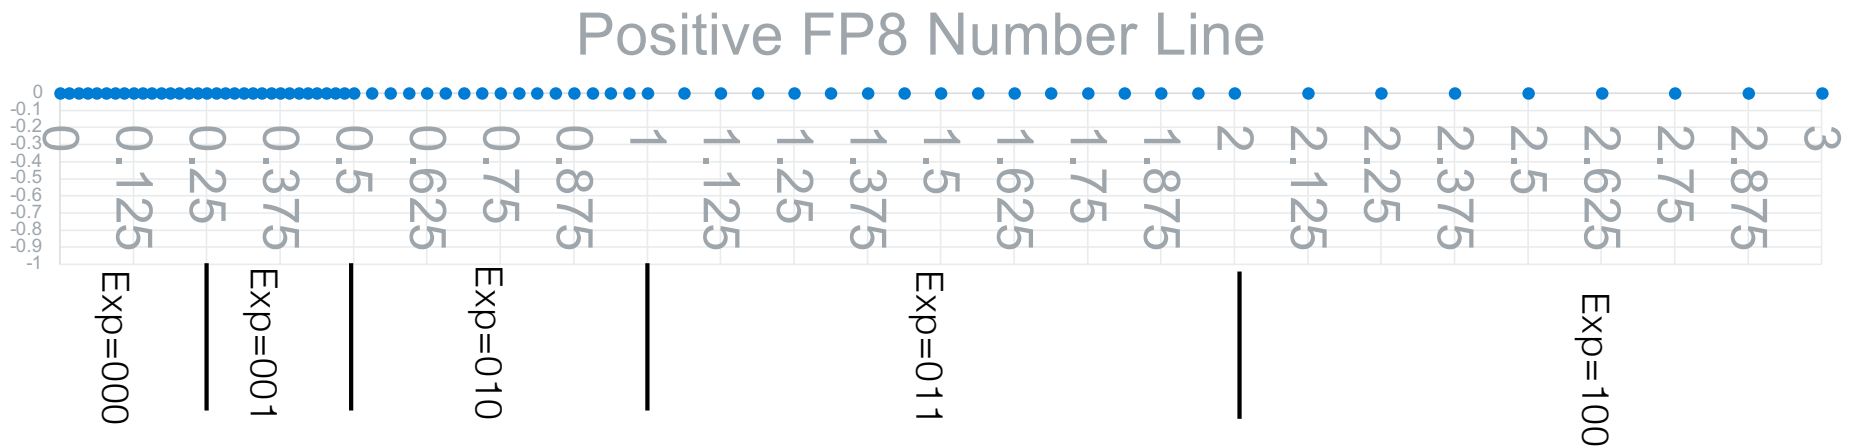
IEE 754 Floating Point

sign	Exp = E + 127	mantissa
1	8	23

sign	Exp = E + 1023	mantissa
1	11	52

- Evolving Standard
- **Single – 32 bit "C Float"**
- **Double – 64 bit "C Double"**
- Half – 16 bit
- Quad – 128 bit
- Binary Floating Point
- Standard also defined decimal FP (not supported by most hardware)
- Special Encodings
 - NAN – not a number (quiet, signaling)
 - $+\infty$ and $-\infty$ (biggest positive #, and smallest negative number)
- Subnormal Numbers

Floating Point and Linearity




Why the non-linearity?
Note, Exp=000 is treated specially

IEEE Floating Point and "C" FP Types

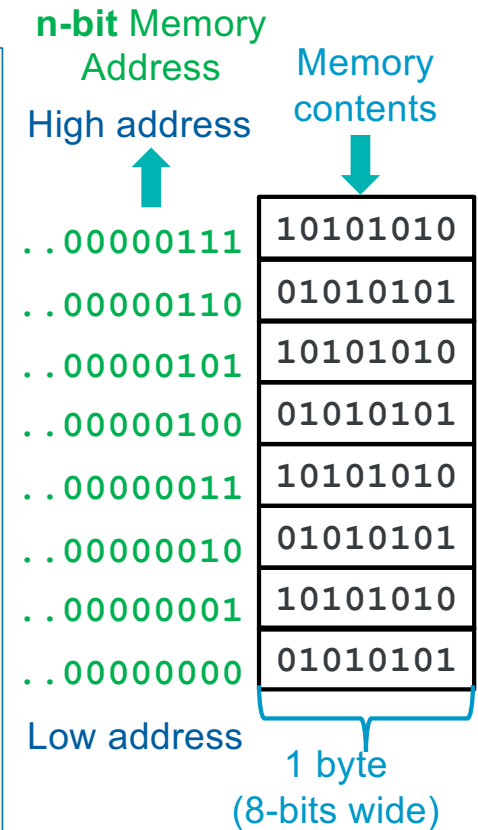
IEEE Type	size	sign	exponent	mantissa	"C" name
bfloat (not IEEE yet)	16	1	8	7	(subset of float)
half	16	1	5	10	
single	32	1	8	23	
double	64	1	11	52	
quad	128	1	15	112	

Memory Review: Organized in Units of Bytes

- One bit (digit) of storage (in memory) has two possible **states**: 0 or 1
- Memory is organized into a **fixed unit** of 8 bits, called a **byte**




- Conceptually, memory is a **single, large array of bytes**, where each **byte** has a unique **address** (*byte addressable memory*)
- An address is an **unsigned** (positive #) *fixed-length* n-bit binary value
 - Range (domain) of possible addresses = **address space**
- Each **byte** in memory can be **individually accessed** and operated on given its **unique address**




Memory Size

- Since memory addresses are implemented in hardware using binary
 - The **Size (number of byte sized cells)** of Memory is specified in **powers of 2**
- Memory size/capacity in **bytes** is specified by the “**Number of bits**” in an address
 - 32 bits of address = $2^{32} = 4,294,967,296$
 - Address Range is 0 to $2^{32} - 1$ (unsigned)
- Shorthand notation for address size (Memory Capacity)
 - KB = 2^{10} (K=1024) kilobyte
 - MB = 2^{20} megabyte
 - GB = 2^{30} gigabyte
 - TB = 2^{40} terabyte
 - PB = 2^{50} petabyte




Variables in Memory: Size and Address

- The number of **contiguous bytes** a variable uses is based on the *type* of the variable
 - Different **variable types** require different numbers of **contiguous bytes**
- **Variable names** map to a starting address in memory
- Example Below: Variables all starting at address 0x80




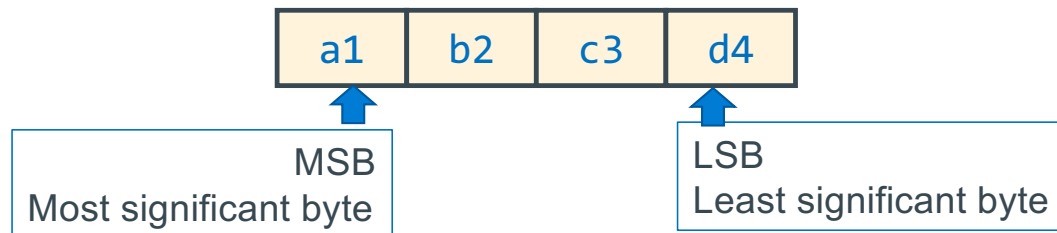
Address and Pointers

- An **address** refers to a location in memory, the **lowest** or **first byte** in a **contiguous sequence of bytes**
- A **pointer** is a **variable** whose **contents** (or value) can be properly used as an **address**
 - The **value in a pointer** *should* be a **valid address allocated to the process by the operating system**
- The **variable x** is at **memory address 0x00001008**
- The **variable pt** is at **memory location 0x00001000**
- The contents of **pt** is the **address of x 0x00001008**



Byte Ordering of Numbers In Memory: Endianness

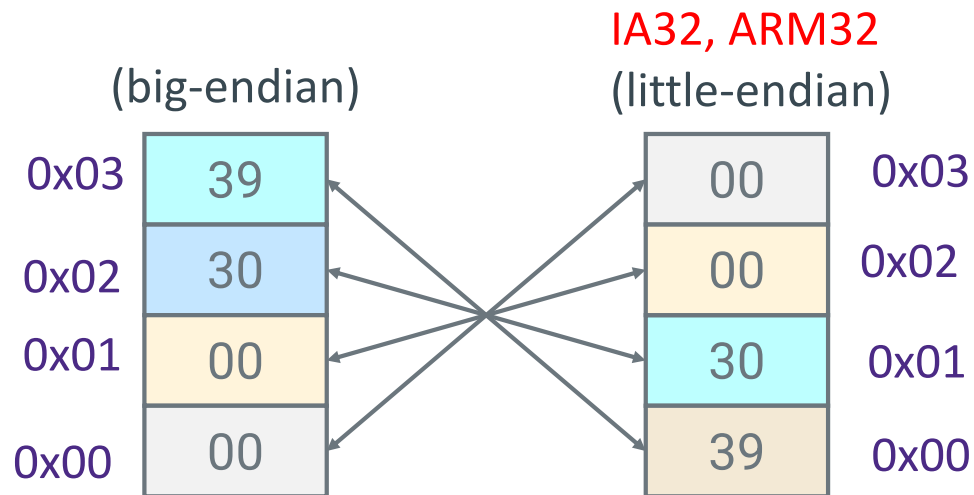
- Two different ways to place multi-byte integers in a **byte addressable** memory
- **Big-endian**: **Most** Significant Byte (“**big end**”) starts at the **lowest (starting)** address
- **Little-endian**: **Least** Significant Byte (“**little end**”) starts at the **lowest (starting)** address
- Example: 32-bit integer with 4-byte data



Byte Ordering Example

Decimal:	12345
Binary:	0011 0000 0011 1001
Hex:	3 0 3 9

```
int x = 12345;  
// or x = 0x3039;
```




Extra Slides

1's Complement Signed Integer Method

- Use the MSB is the sign bit to represent a negative value is encoded as the 1's complement
- All **negative values** have a **one in the leftmost bit**
- All **positive values** have a **zero in the leftmost bit**

Number	+	-
0	0000	1111
1	0001	1110
2	0010	1101
3	0011	1100
4	0100	1011
5	0101	1010
6	0110	1001
7	0111	1000

- The problem is **there are two values for zero**
 - 1111 and 0000 in 4-bits
 - arithmetic is tricky when you cross over the zeros




Negative Integer Numbers: Sign + Magnitude Method



- Use the **Most Significant Bit** as a sign bit
 - 0 as the MSB represents positive numbers
 - 1 as the MSB represents negative numbers
- **Two** (oops) representations for **zero**: 0000, 1000
- Tricky Math (must handle sign bit independently)

$\begin{array}{r} 4 \\ - 3 \\ \hline 1 \end{array}$	$\begin{array}{r} 0100 \\ - 0011 \\ \hline 0001 \end{array}$	$\begin{array}{r} 4 \\ + -3 \\ \hline -7 \end{array}$	$\begin{array}{r} 0100 \\ + 1011 \\ \hline 1111 \end{array}$
	✓		✗

- With Simple math, Positive and Negatives
“increment” (+1) in the **opposite directions!**



Signed Magnitude Examples (Sign bit is always MSB)

0 110
positive 6

1 011
negative 3

Examples (4 bits):

1 000 = -0?	0 000 = 0?
1 001 = -1	0 001 = 1
1 010 = -2	0 010 = 2
1 011 = -3	0 011 = 3
1 100 = -4	0 100 = 4
1 101 = -5	0 101 = 5
1 110 = -6	0 110 = 6
1 111 = -7	0 111 = 7

0 00000000
positive 0

1 0001100
negative 12

Examples Using Hex notation (8 bits):

0x00 = 0b00000000 is positive, because the sign bit is 0

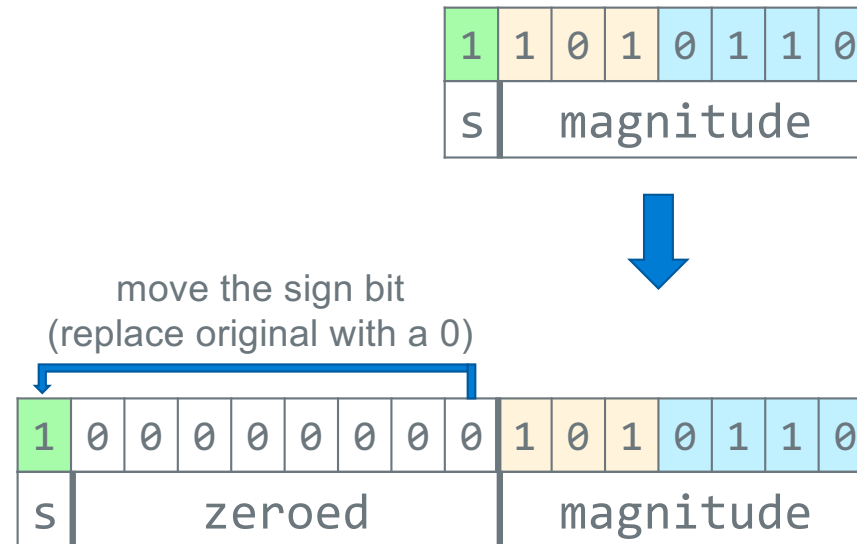
0x7F = 0b01111111 is positive (+127₁₀)

0x85 = 0b10000101 is negative (-5₁₀)

0x80 = 0b10000000 is negative... also zero

Sign Extension Signed Magnitude number

- Just move the sig bit and expand the magnitude with zeros to the left




Interpreting and extending with Different representations

How to extend this bit pattern?

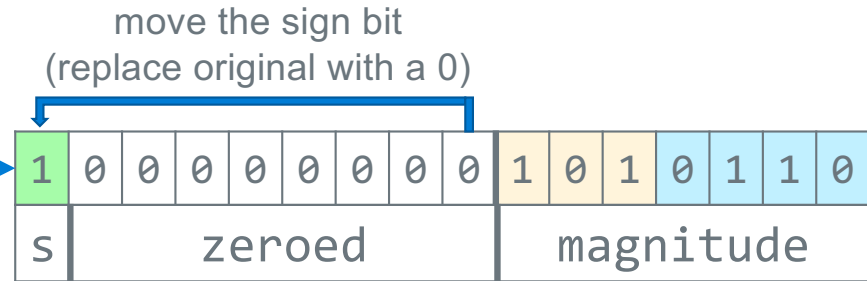
0xd6

unsigned



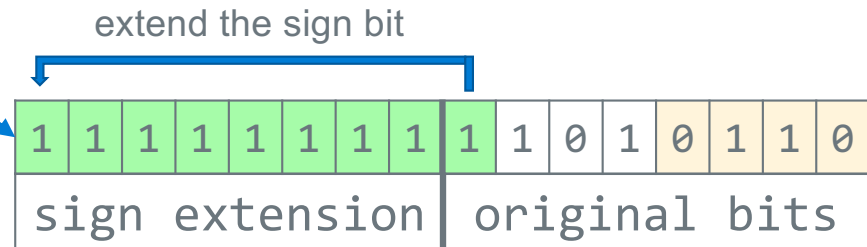
0x00d6

signed
magnitude



0x8056

two's
complement



0xffd6

Watch out for Hardware differences: Example: Is a Char signed or unsigned?

```
#include <stdio.h>
#include <stdlib.h>

int
main(void)
{
    char c = 255;

    printf("%d\n", (int)c);

    return EXIT_SUCCESS;
}
```

- variable c is being cast promoted to an int
- So, what is printed?
- Depends on the hardware
- Aarch32 and Aarch64 (arm) it is unsigned
255
- Intel 64-bit it is signed
-1

sizeof(): Variable Size (number of bytes) Operator

```
#include <stddef.h>
/* size_t type may vary by system but is always unsigned */
```

sizeof() operator returns:

the number of bytes used to store a variable or variable type

```
size_t size = sizeof(variable_type);
```

or

```
size_t size = sizeof(variable_name); // preferred!
```

- The argument to sizeof() is often an expression:

```
size = sizeof(int * 10);
```

- reads as:

- number of bytes required to store **10 integers (an array of [10])**


sizeof() Examples (On the pi-cluster)

```
#include <stdio.h>
#include <stdlib.h>
int
main (void){

    printf("char   is %zu\n", sizeof(char));
    printf("short  is %zu\n", sizeof(short));
    printf("int     is %zu\n", sizeof(int));
    printf("long    is %zu\n", sizeof(long));
    return EXIT_SUCCESS;
}
```

```
% ./a.out (on the picluster)
char   is 1
short  is 2
int     is 4
long    is 4
```

Floating Point Number in a Byte (Not A Real Format)



- **Mantissa encoding:** = 1.[xxxx] encoded as an unsigned value
- **Exponent encoding:** 3 bits encoded as an unsigned value using bias encoding
 - Bias encoding = $(2^{E-1} - 1)$
 - 3 bits for the bias we have $2^{3-1} - 1 = 2^2 - 1 =$ a bias of 3
 - **With a Bias of 3:** positive and negative numbers range: small to large is: 2^{-3} to 2^4

Actual	-3	-2	-1	0	1	2	3	4
Bias	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+3
Biased	0	1	2	3	4	5	6	7

Encoding Fractions Observations

Examples In Base 10:

$$42.4 \times 10^5 = 4.24 \times 10^6$$

$$324.5 \times 10^5 = 3.245 \times 10^7$$


$$0.624 \times 10^5 = 6.24 \times 10^4$$

Observation on base 10:

We usually adjust the exponent until we get down to one digit to the left of the decimal point

Characteristics of Signed Numbers

- Digital Hardware (and C) supports two flavors of integers
 - *unsigned* – only non-negative (positive) numbers
 - *signed* – both negative and non-negatives (positive) numbers
- A Signed integer must be able to represent:
 - Negative integer
 - Zero (0)
 - Positive Integer
 - $\text{number} + (- \text{representation of number}) = 0$
- So, with a fixed number of bits, some of the bit patterns in the wheel at the left must be reallocated to represent negative numbers



Example: Two's Complement Sign or bit Extension - 3

- Adding 1's if front of a negative number does not change its value

$$\begin{array}{rcl} 1 & = & 0001 \\ & & \downarrow \downarrow \downarrow \downarrow \\ \text{invert} & = & 1110 \\ \text{add } 1 & + & \quad 1 \\ -1 & = & \underline{1111} \end{array}$$

$$\begin{array}{rcl} -1 & = & 11111111 \\ \text{extend to} & & \\ \text{16 bits} & & \end{array}$$

1111111111111111

$$\begin{array}{rcl} -1 & = & 1111 \\ \text{extend to} & & \\ \text{8 bits} & & \end{array}$$

11111111

Sign Extension in C: Type casts

- Convert from smaller to larger integral data types
- C and Java automatically performs sign extension
- Example (remember we are working with 32-bit int and 16-bit short)

0b0011

```
short int sx = 12345;
int      ix = (int) sx;
```

Var	Decimal	Hex	Binary
sx	12345	30 39	00110000 00111001
ix	12345	00 00 30 39	00000000 00000000 00110000 00111001

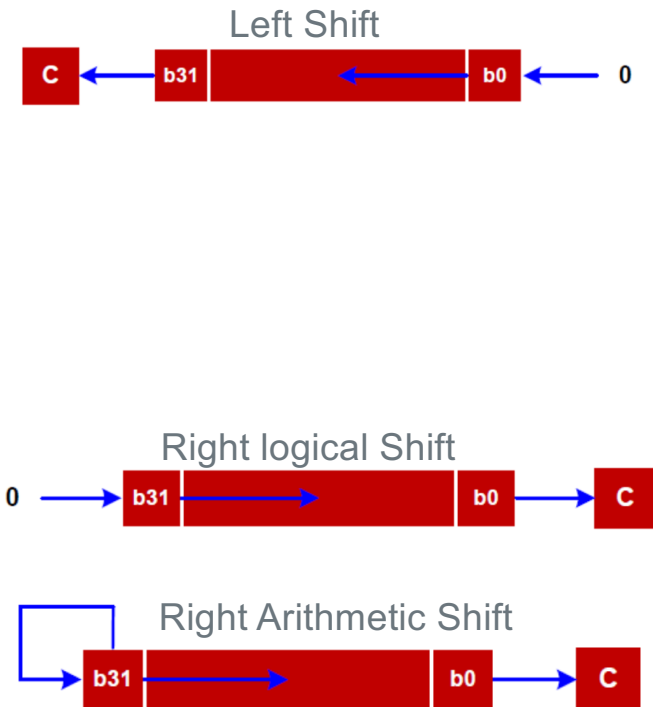
0b1100

```
short int sy = -12345;
int      iy = (int) sy;
```

Var	Decimal	Hex	Binary
sy	-12345	CF C7	11001111 11000111
iy	-12345	FF FF CF C7	11111111 11111111 11001111 11000111

Shift Operations in C

- n is number of bits to shift a variable x of width w bits
- Shifts by $n < 0$ or $n \geq w$ are *undefined*
- Left shift ($x \ll N$)
 - Shift N bits left, Fill with 0s on right
- In C: behavior of \gg is determined by compiler
 - gcc: it depends on data type of x (signed/unsigned)
- Right shift ($x \gg N$)
 - Logical shift (for unsigned variables)
 - Shift N bits right, Fill with 0s on left
 - Arithmetic shift (for signed variables) – Sign Extension
 - Shift N bits right while Replicating the most significant bit on left
 - Maintains sign of x
- In Java: logical shift is \ggg and arithmetic shift is \ggg




Arithmetic Shift Right 1 Digit = Divide by 2 for 2's Complement Values



Most significant Digit Initial 8-digit Value Least significant digit



Binary
Positive Number



Replicate MSB

Carry Bit
"bit bucket"

Final 8-digit Value

Number ranges: limits.h

- The include file `<limits.h>` defines various symbolic names where the names represent the various limits on resources that the implementation imposes on applications;

C Data Type	signed min	signed max	unsigned max
char	SCHAR_MIN	SCHAR_MAX	UCHAR_MAX
short int	SHRT_MIN	SHRT_MAX	USHRT_MAX
int	INT_MIN	INT_MAX	UINT_MAX
long int	LONG_MIN	LONG_MAX	ULONG_MAX

```
#include <limits.h>
int i = INT_MAX;
...
    printf("Max int is: %d\n", i);
...
```

- The standard C integer types were intended to allow code to be portable among machines with different inherent data sizes (word sizes), so each type may have different range of numbers on different machines

Excess N Bias Encoding Method

- **Excess, Bias (or offset) encoding** maps negative numbers to an unsigned (positive) integer range by adding an offset number (called the bias) to encode positive and negative numbers
 - **Most negative number maps to zero**, **most positive number maps to all 1's**
- **For example: Say we have a number that is limited to 3 bits (0 to 7 unsigned)**

Actual	-3	-2	-1	0	1	2	3	4
Bias	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+3
Biased Encoded	0	1	2	3	4	5	6	7

Excess Bias Encoding (As used in floating point numbers)

- Given a number in E bits, to divide the range in about $1/2$ the following is used:

$$\text{excess N bias} = (2^{E-1} - 1) \quad (\text{this is just one of many bias formulas})$$

- With this excess N Bias approach:** actual numbers range from most negative to most positive is: **-(bias) to bias+1**
- So, for a number that is limited to 4 bits (0 to 15 unsigned)**
 - Then excess N bias = $2^{4-1} - 1 = 2^3 - 1 =$ a bias of +7

actual	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6	7	8
bias	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7	+7
bias encoded	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15

Another Way to Look at 2's Complement Encoding

- A 2's complement value can be thought of as using a slightly different **bias encoding** for negative numbers only (more negative values): -2^{w-1}
- The **leftmost bit** is then interpreted as a **decision to apply the bias** (if **1**) or not (if **0**)
 - **1** apply the bias
 - **0** do not apply the bias
- For example, for a 4-bit number ($w = 4$), the negative number bias weight would be $= -2^{4-1} = -2^3 = -8$

2's	1000	1001	1010	1011	1100	1101	1110	1111	0000	0001	0010	0011	0100	0101	0110	0111
3 bit	000	001	010	011	100	101	110	111	000	001	010	011	100	101	110	111
decimal	0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7
+Bias	-8	-8	-8	-8	-8	-8	-8	-8	0	0	0	0	0	0	0	0
Actual	-8	-7	-6	-5	-4	-3	-2	-1	0	1	2	3	4	5	6	7

Observe: adding +1 makes the number more positive for both negative and positive numbers

Scientific Notation Decimal

The diagram illustrates the components of the scientific notation expression $+6.02_{10} \times 10^{23}$. Labels with arrows point to specific parts: 'sign' points to the '+' sign; 'mantissa' points to the '6.02' part with a bracket above it; 'decimal point' points to the '.' in '6.02'; 'exponent' points to the '23' in the power of 10; and 'radix (base)' points to the '10' in the power of 10.

- *Scientific Normalized form:*


exactly one digit (non-zero) to left of decimal point

- Alternatives to representing 1/1,000,000,000

- **Normalized:** 1.0×10^{-9}

- Not normalized: 0.1×10^{-8} , 10.0×10^{-10}

Floating Point Number in a Byte (Not A Real Format)



- **Mantissa encoding:** = 1.[xxxx] encoded as an unsigned value
- **Exponent encoding:** 3 bits encoded as an unsigned value using bias encoding
 - Bias encoding = $(2^{E-1} - 1)$
 - 3 bits for the bias we have $2^{3-1} - 1 = 2^2 - 1 =$ a bias of 3
 - **With a Bias of 3:** positive and negative numbers range: small to large is: 2^{-3} to 2^4

Actual	-3	-2	-1	0	1	2	3	4
Bias	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+ 3	+3
Biased	0	1	2	3	4	5	6	7

Floating Point Number (8-bits) Number Range: 2^{-3} to 2^4




0.0 Special case in this simple model
we do not put back the “hidden bit”



Smallest Non-zero Positive
 $0.00\textcolor{blue}{1}0001 = \textcolor{blue}{1}/8 + 1/128 = 0.1328125$ base 10




Largest Positive/Negative
 $\textcolor{blue}{1}.\textcolor{blue}{1111} \times 2^4 = \textcolor{blue}{11111} = 31$ base 10



Smallest (closest to zero) Number
 $\textcolor{blue}{1}.\textcolor{blue}{0000} \times 2^{-3} = 0.00\textcolor{blue}{1}000 = \textcolor{blue}{1}/8 = -0.125$ base 10

Note: Orange is hidden bit added back

Decimal to Float



Step 1: convert from base 10 to binary (absolute value)

$$-0.375 (\text{decimal}) = 0000.0110_2$$

Step 2: Find out how many places to shift to get the number into the normalized 1.xxxx mantissa format

$$0000.0110_2 = 1.1000 \times (2^{-2})_{\text{base } 10}$$

$$\text{exponent: } -2_{10} + \text{bias of } 3_{10} = 1_{10} = 0b001 \text{ for the exponent (after adding the bias)}$$

Step 3: Use as many digits that fit to the right of the decimal point in the fractional .xxxx part

$$1.1000$$


Step 4: Sign bit

positive sign bit is 0

negative sign bit is 1

s	exponent	fraction
1	0b001	0b1000
0x9		0x8

$$= 0x98$$




Step 1: Break into binary fields

0x45 =

	0x4	0x5
s	exponent	fraction
0	0b100	0b0101

Step 2: Extract the unbiased exponent


0b100 = $4_{\text{base } 10} - \text{bias of } 3_{10} = 1_{10}$ for the exponent (bias removed)

Step 3: Express the mantissa (restore the **hidden** bit)

1.0101

Step 4: Apply the **unbiased** exponent

1.0101_{base 2} × $(2^1)_{\text{base } 10} = \mathbf{10.101}$


Step 5: Convert to decimal

10.101 = $2.625_{\text{base } 10}$

Step 6: Apply the Sign


+ $2.625_{\text{base } 10}$

IEEE “754” Floating Point Double and Single Precision



Bias is $(2^{8-1} - 1) = 127$


single precision floating point number = $(-1)^s \times 2^{E-127} \times 1.\text{fraction}$



bias is $(2^{11-1} - 1) = 1023$

double precision floating point number = $(-1)^s \times 2^{E-1023} \times 1.\text{fraction}$

Decimal to IEEE Single Precision Float



Step 1: convert from base 10 to binary (absolute value)

$$-13.375(\text{decimal}) = 1101.0110$$

Step 2: Find out how many places to shift to get the number into the normalized 1.xxxx mantissa format


$$1101.0110 = 1.1010110 \times (2^3)_{\text{base } 10}$$

$$3 + \text{bias of } 127 = 130 \text{ for the exponent} = 0b1000 \ 0010$$


Step 3: Use as many digits that fit to the right of the decimal point in the fractional .xxxx part (0 pad)

$$1.1010110 \ 0000 \ 0000 \ 0000 \ 0000$$

Step 4: If the sign is positive sign bit is 0, otherwise it is 1




IEEE Single Precision Float to Decimal



Step 1: Break into binary fields and **expand** as needed

$0xc0b00000 =$

Step 2: Find the exponent



$0b10000001 = 129_{\text{base } 10} - \text{bias of } 127_{10} = 2_{10}$ exponent with bias added

Step 3: Express the mantissa (restore the hidden bit)

1.0110

Step 4: Apply the exponent

$1.0110 \times (2^2)_{\text{base } 10} = 101.10$

Step 5: Convert to decimal

$101.10 = 5.5$

Step 6: Apply the Sign

-5.5

Reference: 8-Bit Overflow Examples

Unsigned Integer

cout										
1	1	1	1	1	1	1	0		carries	
	1	1	1	0	1	0	1	0		=234 ₁₀
+	0	0	1	1	0	1	1	0		=54 ₁₀
	0	0	1	0	0	0	0	0		=32 ₁₀

Because carry-out bit is 1 (and dropped), overflow is detected

Two's Complement

cout	cin									
0	1	1	1	1	1	1	0		carries	
	0	1	1	0	1	0	1	0		=106 ₁₀
+	0	0	1	1	0	1	1	0		=54 ₁₀
	1	0	1	0	0	0	0	0		=-96 ₁₀


Both operands are positive, but resulting sign is negative
see that cout != cin at the MSB
overflow is detected

Two's Complement

cout	cin									
1	1	1	1	1	1	1	0		carries	
	1	1	1	0	1	0	1	0		=-22 ₁₀
+	0	0	1	1	0	1	1	0		=54 ₁₀
	0	0	1	0	0	0	0	0		=32 ₁₀

Unlike unsigned arithmetic, no overflow even though the carry-out bit is 1.
As the operand's signs differ, overflow is not possible (cout == cin)

Decimal to Float



Step 1: convert from base 10 to binary (absolute value)

$$6.625 (\text{decimal}) = 0110.1010$$

Step 2: Find out how many places to shift to get the number into the normalized 1.xxxx mantissa format

$$0110.1010 \text{ normalizes to } \rightarrow 1.101010 \times (2^2)_{\text{base } 10}$$

$$\text{exponent: } 2_{10} + \text{a bias of } 3_{10} = 5_{10} = 0b101 \text{ for the exponent (after adding the bias)}$$

Step 3: Use as many digits to the right of the decimal point that will fit in the fractional .xxxx part

$$1.1010\underline{10} \text{ (we will truncate drop the trailing } \underline{10}, \text{ Real FP use complex rounding approaches)}$$

Step 4: Sign bit

positive sign bit is 0

negative sign bit is 1

s	exponent	fraction
0	0b101	0b1010
0x5		0xa

$$= 0x5a$$