

Version 1.05


UCSD CSE 30

Computer Organization and Systems Programming

C Programming Part 2

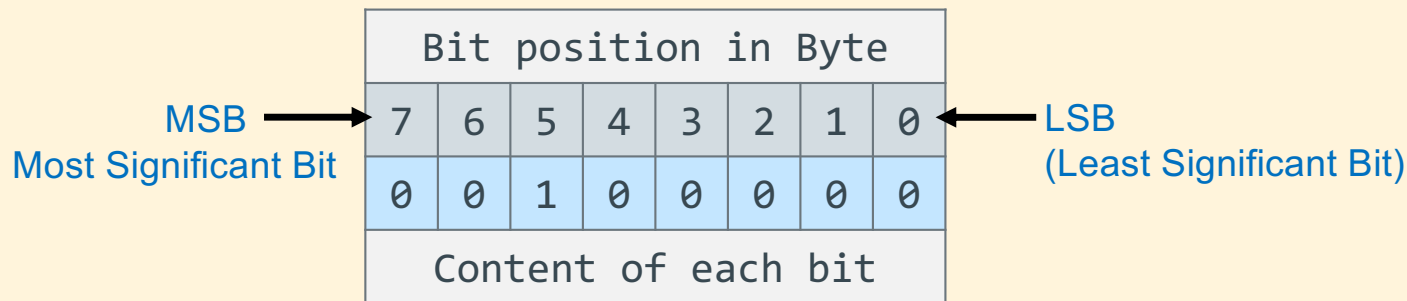
Lecture 6 – Oct 11, 2022

Keith Muller

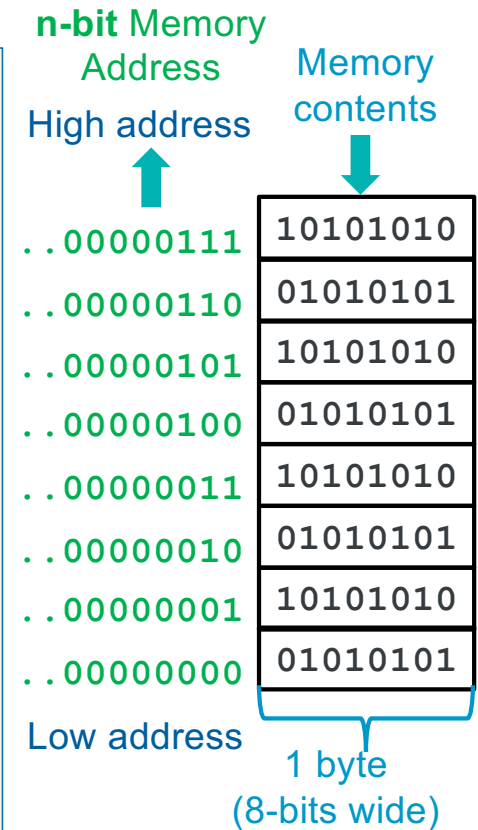


Memory Review: Organized in Units of Bytes

- One bit (digit) of storage (in memory) has two possible **states**: 0 or 1
- Memory is organized into a **fixed unit** of 8 bits, called a **byte**

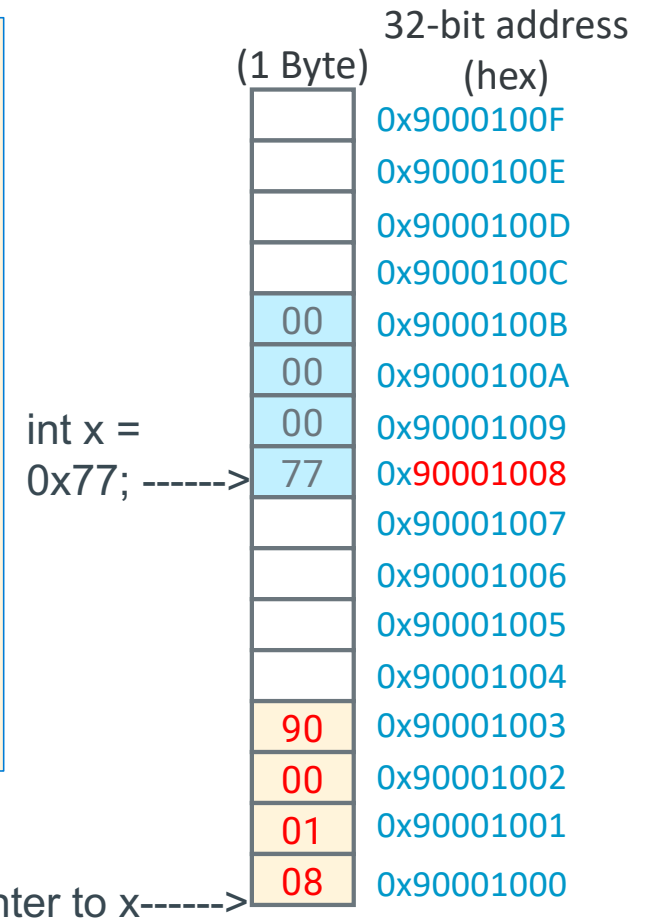


- Conceptually, memory is a **single, large array of bytes**, where each **byte** has a unique **address** (*byte addressable memory*)
- An address is an **unsigned** (positive #) *fixed-length* n-bit binary value
 - Range (domain) of possible addresses = **address space**
- Each **byte** in memory can be **individually accessed** and operated on given its **unique address**



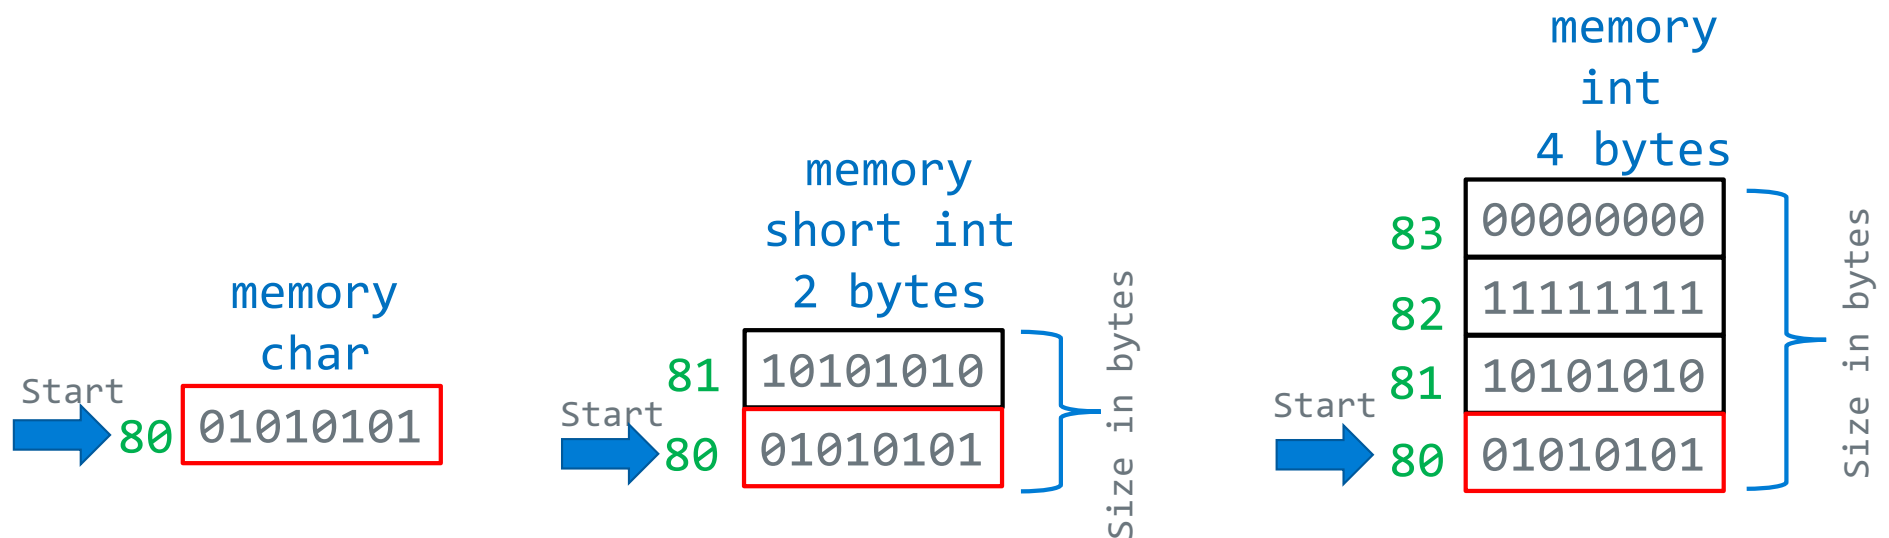
Address and Pointers

- An **address** refers to a location in memory, the **lowest** or **first byte** in a **contiguous sequence of bytes**
- A **pointer** is a **variable** whose **contents** (or value) can be properly used as an **address**
 - The **value in a pointer** *should* be a **valid address allocated to the process by the operating system**
- The **variable x** is at **memory address 0x90001008**
- The **variable pt** is at **memory location 0x90001000**
- The **contents** of **pt** is the **address of x 0x90001008**



Variables in Memory: Size and Address

- The **number of contiguous bytes** a variable uses is based on the *type* of the variable
 - Different **variable types** require different numbers of **contiguous bytes**
- **Variable names** map to a starting address in memory
- **Example Below:** Variables all starting at address 0x80, each box is a byte



sizeof(): Variable Size (number of bytes) Operator

```
#include <stddef.h>
/* size_t type may vary by system but is always unsigned */
```

sizeof() operator returns a value of type **size_t**:

the number of bytes used to store a variable or variable type

```
size_t size = sizeof(variable_type);
```

or

```
size_t size = sizeof(variable_name); // preferred!
```

- The argument to sizeof() is often an expression:

```
size = sizeof(int * 10);
```

- reads as:


- number of bytes required to store **10 integers (an array of [10])**

Memory Addresses & Memory Content

```
x = x;    // Lvalue = Rvalue
```

Variable name in a C statement evaluates to either:


- **Lvalue:** when on the **left side** (Lside or Left value) of the **=** sign is the
 - address where it is stored in memory – a constant
 - Address assigned to a variable cannot be changed at runtime
- **Rvalue:** when on the **right side** (Rside or Right value) of an **=** sign is the
 - contents or value stored in the variable (at its memory address)
 - requires a memory read to obtain



Memory Addresses & Memory Content

```
y = 42;
```

```
x = y;    // Lvalue = Rvalue
```



- **x** on left side (**Lside**) of the assignment operator = evaluates to:
 - The address of the memory assigned to the **x** – this is x's **Lvalue**
- **y** on right side (**Rside**) of the assignment operator = evaluates to:
 - **READ** the contents of the memory assigned to the variable **y** (type determines length – number of bytes) - this is y's **Rvalue**
- So `x = y;` is:

Read memory at **y** (**Rvalue**); write it to memory at **x's address** (**Lvalue**)

Introduction: Address Operator: &

- Unary **address operator** (&) produces the **address** of where an **identifier** is in memory
- Requirement: **identifier must have a Lvalue**
 - Cannot be used with **constants** (e.g., 12) or **expressions** (e.g., x + y)
 - **&12** does not have an **Lvalue**, so **&12** is not a legal expression
- How can I get an **address for use on the Rside**? Three ways:
 - **&var** (any variable identifier or name)
 - **function_name** (name of a **function**, not func()); **&funct_name** is equivalent
 - **array_name** (name of the **array** like array_name[5]); **&array_name** is equivalent

Introduction: Address Operator: &

- Unary **address operator** (&) produces the **address** of where an **identifier** is in memory

- **Example:** this might print:
value of g is: 42
address of g is: 0x71a0a0
(the address will vary)

```
int g = 42;
int
main(void)
{
    printf("value of g is: %d\n", g);
    printf("address of g is: %p\n", &g);
    return EXIT_SUCCESS;
}
```

- **Tip:** printf() format specifier to display an address/pointer (in hex) is "%p"

Introduction: Pointer Variables - 1

- In C, there is a *variable type* for storing an address: a *pointer*
 - **Contents** of a pointer is an unsigned (0+, positive numbers) memory address
- When the **Rside** of a variable contains a *memory address*, (it *evaluates* to an *address*) the variable is called a *pointer variable*
- A *pointer* is defined by placing a *star* (or *asterisk*) (*) before the identifier (name)

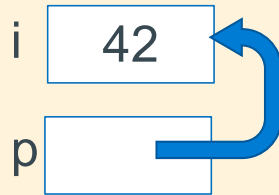
```
type *name; // defines a pointer; name contains address of a variable of type
```

Introduction: Pointer Variables - 1

`type *name; // defines a pointer; name contains address of a variable of type`

- You also must specify the **type of variable** to which the pointer points

```
int i = 42;  
int *p; /* p contains the address of an integer */  
p = &i; /* p "points at" i (assign address of i to p) */
```



- Recommended:** be careful when defining multiple pointers on the same line:

`int *p1, p2;` is not the same as: `int *p1, *p2;`

Use instead:

```
int *p1;  
int *p2;
```

Introduction: Pointer Variables - 2

- Pointers are typed! Why?

- The compiler needs the **size** (`sizeof()`) of the data **you are pointing at** (number of bytes to access)

- A pointer definition:

```
int *p = &i;  /* p points at i (assign address i to p) */
```

- Is the same as writing the following definition and assignment statements

```
int *p;      /* p is defined (not initialized) */  
p = &i;      /* p points at i (assign address i to p) */
```

- The ***** is part of the definition of **p** and is **not part of the variable name**
 - The name of the variable is **simply p**, not ***p**
- C mostly ignores whitespace, so these three definitions are equivalent

```
int  *p = &i;    /* Style A */  
int *p = &i;    /* Style B */  
int*p = &i;    /* Style C */
```

Introduction: Pointer Variables - 3

- As with any variable, its value can be changed

`p = &j;` */* p now points at j */*

i 42

j 77

p



`p = &i;` */* p now points at i */*



i 42

j 77

p



Introduction: Pointer Variables - 4

- Pointer variables all use the **same amount of memory** no matter what they point at


```
int *iptr;  
char *cptr;  
  
printf("iptr(%u) cptr(%u)\n", sizeof(iptr), sizeof(cptr));
```

- Above prints on a 32-raspberry pi `iptr(4) cptr(4)`

Introduction: Indirection (or dereference) Operator: *

- The **indirection operator** (*) or the *dereference operator to a variable* is the **inverse** of the *address operator* (&)
- **address operator** (&) can be thought of as:

“get the address of this box”



- **indirection operator** (*) can be thought of as:

“follow the arrow to the next box and get its contents”

Introduction: Indirection (or dereference) Operator: *

*Contents of **p** is the address of **i** (p points at i)*


```
int i = 42;  
int j = i;  
int *p;  
p = &i;  
  
printf("*p is %d\n", *p);
```

```
% ./a.out  
*p is 42
```

Introduction: Indirection Operator Rside

- Performs the following steps when the ***** is on the Rside:
 1. **read** the **contents** of the **variable** to get an **address**
 2. **read** and return the **contents** at that **address**
 - (requires **two reads of memory on the Rside**)

```
z = *x; // copy the contents of memory pointed at by x to z
```




Introduction: Indirection Operator Lside

Performs the following steps when the ***** is on the Lside:

1. **read** the **contents** of the **variable** to get **an address**
2. **write** the evaluation of the Rside expression to that address
 - (requires **one read of memory and one write of memory on the Lside**)

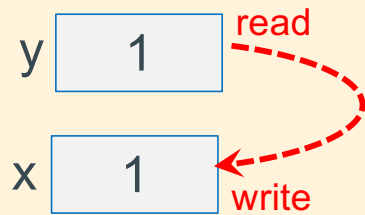
```
*z = x; // copy the value of x to the memory pointed at by z
```



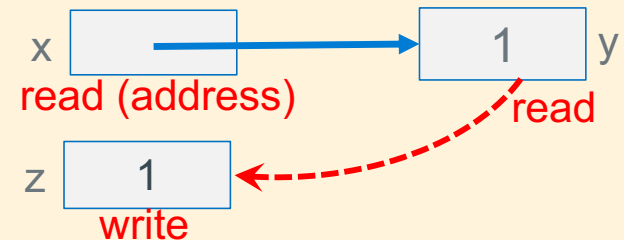
Each use of a * operator results in one additional read -1

Each * when used as a dereference operator in a **statement** (Lside and Rside) generates an additional read

```
int x = 2, y = 1;  
x = y; // one read
```



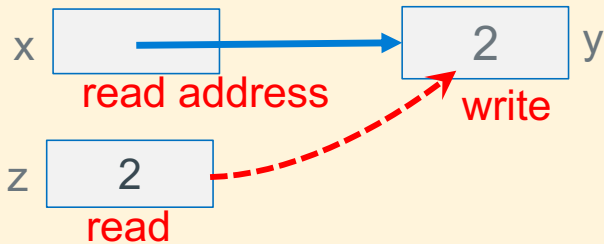
```
int z = 2, y = 1;  
int *x;  
x = &y;  
z = *x; // two reads
```



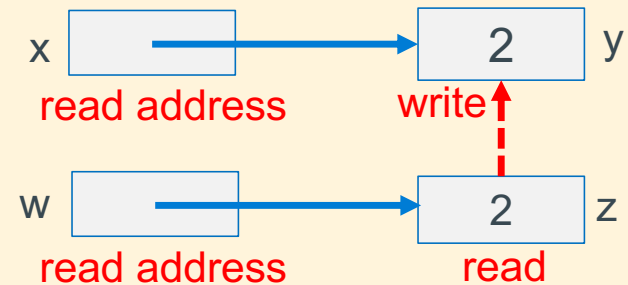
Each use of a * operator results in one additional read -2

- Each * when used as a dereference operator in a **statement** (Lside and Rside) generates an additional read

```
int z = 2, y = 1;  
int *x;  
x = &y;  
*x = z;
```



```
int z = 2, y = 1;  
int *x;  
int *w;  
x = &y;  
w = &z;  
*x = *w;
```



Recap: Lside, Rside, Lvalue, Rvalue

```
int x = 2, y = 1;
x = y;
```


```
int z = 2, y = 1;
int *x;
int *w;
x = &y;
w = &z;
*x = *w;
```

```
*x on Lside is 0x108
w on Rside is 0x100
*w on Rside is 2
```

Constant
Var Name
y
x

Lvalue
address
0x108
0x104

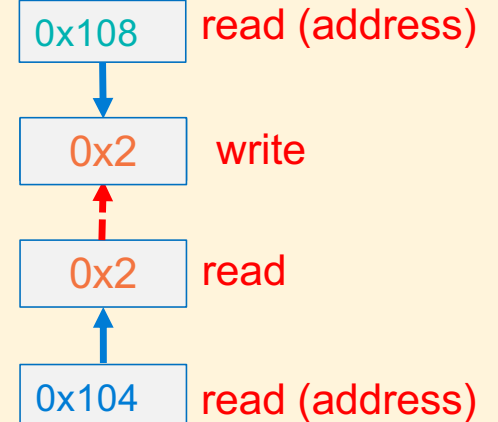
Rvalue
Contents



Constant
Var Name
x
y
z
w

Lvalue
address
0x10c
0x108
0x104
0x100

Rvalue
Contents



Pointer Practice

```
int *ptr;
```

Declares a variable, `ptr`, which is a pointer to (it contains the address of) an `int` in memory

```
int x = 5;
```

```
int y = 2;
```

Declares two variables, `x` and `y`, that contain `ints`, and *initializes* them to 5 and 2, respectively

```
ptr = &x;
```

Sets `ptr` to contain the address of `x` ("`ptr` points to `x`")

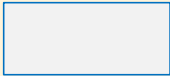
```
y = 1 + *ptr;
```

"Dereference `ptr`"

Sets `y` to "1 plus the value stored at the address held by `ptr`. Because `ptr` points to `x`, this is equivalent to `y = 1 + x`;




```
x = *(&y);
```

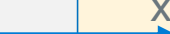


Sets `x = y`; The `*` and `&` cancel each other. get the address of `y` and then get the contents pointed by that address




`ptr` 

`x`  write

`y`  write

`ptr`  `x`  5
`y`  2

`ptr`  `x`  5 read
`y`  6 write

`ptr`  `x`  6 write
`y`  6 read

The NULL Constant and Pointers

- **NULL is a constant** that **evaluates to zero (0)**
- You **assign a pointer variable to contain NULL** to **indicate that the pointer does not point at anything**
- A **pointer variable** with a **value of NULL** is called a “**NULL pointer**” (invalid address!)
- Memory location 0 (address is 0) is not a valid memory address in any C program
- Dereferencing NULL at runtime will cause a program fault (segmentation fault)!

```
p = NULL;  
i = *p;           /* segmentation fault! */  
*(int *)900000 = 25; /* cast 900000 to a pointer */  
                  /* if writeable address space, it works */  
                  /* that memory location just changed */
```

Using the NULL Pointer

- Many functions return NULL to indicate an error has occurred

```
/* these are all equivalent */  
int *p = NULL;  
int *p = (int *)0;    // cast 0 to a pointer type  
int *p = (void *)0;   // automatically gets converted to the correct type
```

- NULL is considered “false” when used in a Boolean context
 - Remember: false expressions** in C are defined to be zero or NULL
- The following two are equivalent (the second one is preferred for readability):


```
if (p) ...  
if (p != NULL) ...
```

What is Aliasing?

- **Two or more** variables are **aliases** of each other when they all **reference the same memory** (so different names, same memory location)
- When one pointer is copied to another pointer it **creates an alias**
- **Side effect**: Changing one variables value (content) changes the value for other variables
 - **Multiple variables** all **read and write** the **same** memory location
 - Aliases occur either by **accident** (coding errors) or **deliberate** (**careful: readability**)

```
int i = 5;  
int *p;  
int *q;  
p = &i;
```

```
q = p;    // *p & *q are aliases  
*q = 4;   // changes i
```



Result *p, *q and i all have the value of 4

Version 1.05


UCSD CSE 30

Computer Organization and Systems Programming

C Programming Part 2

Lecture 7 – Oct 13, 2022

Keith Muller



Defining Arrays

Definition: **type** name[**count**]

- **"Compound"** data type where each value in an array is an element of **type**
- Allocates **name** with a **fixed count** array elements of type **type**
- Allocates (**count** * **sizeof(type)**) bytes of **contiguous memory**
- Common usage is to specify a compile-time constant for **count**


```
#define BSZ 6  
int b[BSZ];
```

BSZ is a macro replaced by the C preprocessor at compile time

- Array **names are constants (like all variable names)** and **cannot be assigned** (the name cannot appear on the Lside by itself)

```
a = b;           // invalid does not copy the array  
                // copy arrays element by element
```

```
int b[6];
```



Array Initialization


- Initialization: `type name[count] = {val0,...,valN};`
 - `{ }` (*optional*) initialization list can only be used at **time of definition**
 - If no `count` supplied, `count` is determined by compiler using the number of array initializers

no initialization values given; then elements are initialized to 0
 - `int block[20] = {};` *//only works with constant size arrays*
 - defines an **array of 20 integers** each element filled with zeros
 - Performance comment: do not zero automatic arrays unless really needed!
 - When a **count** is given:
 - extra initialization values** are **ignored**
 - missing initialization values** are set to **zero**

```
int block[5] = {2, 3, 5, 6, 11, 13};
```

not needed and if used **may** truncate initialization list

6 initialization values given, **only 5 are used**



X

Accessing Arrays Using Indexing


- **name** [**index**] selects the **index** element of the array
 - **index should be unsigned**
 - **Elements range from: 0 to count – 1** (`int x[count];`)
- **name** [**index**] can be used as an **assignment target** or as a **value in an expression**
- **Array name** (by itself with no []) on the **Rside** evaluates to the **address of the first element of the array**

```
int a[5];  
int b[5];
```

```
int b[5];  
int *p = b;
```

p 9000

1 word
(int = 4 bytes)




How many elements are in an array?

- The number of elements of space allocated to an array (called **element count**) and indirectly the total size in bytes of an array **is not stored anywhere!!!!!!**
- An **array name** is just the **address of the first element in a block of contiguous memory**
 - So an array does not know its own size!

```
#define SZ 6
int block[SZ];    // you specify the array has SZ elements
int indx;        // use when SZ is defined

for (indx = 0; indx < SZ; indx++)
    block[indx] = 0;
```

```
int b[6];
```



Determining Element Count for a compiler calculated array

- Programmatically determining the element count in a compiler calculated array
`sizeof(array) / sizeof(of just one element in the array)`
- `sizeof(array)` **only works** when used in the SAME **scope** as where the array variable was **defined**

```
#include <stddef.h>

int block[] = {2, 3, 5, 6, 11, 13};    // automatic: compiler calculates array size

int cnt = (int)(sizeof(block) / sizeof(block[0])); // in this case cnt = 6


for (int indx = 0; indx < cnt; indx++)
    block[indx] = 0;
```

Pointer and Arrays - 1

- A few slides back we stated: **Array name** (by itself) on the **Rside** evaluates to the **address of the first element of the array**
- Array indexing syntax (`[]`) an operator that performs *pointer arithmetic*
- `buf`** and **`&buf[0]`** on the **Rside** are **equivalent**, **both evaluate** to the address of the first array element

```
int *p = buf;           // or int *p = &buf[0];
int *p1 = &buf[1];
int *p2 = &buf[2];
int *p3 = &buf[3];

*p = *p + 10;
*p1 = *p1 + 10;         // {12, 13, 5, 6, 11}
```



Pointer and Arrays - 2

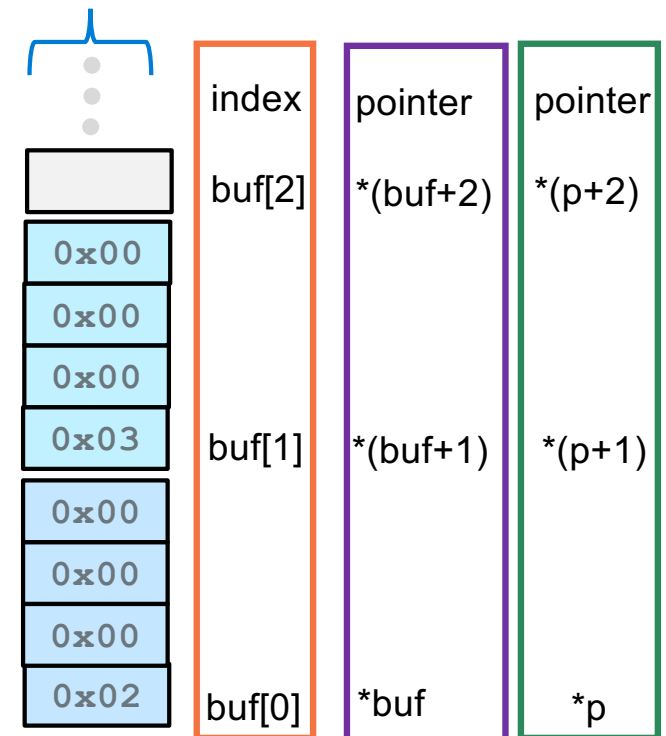
When `p` is a pointer, the actual value of `(p+1)` depends on the type that pointer `p` points at

- `(p+1)` adds `1 x sizeof(what p points at)` bytes to `p`
 - `++p` is equivalent to `p = p + 1`
- Using pointer arithmetic to find array elements:
 - Address of the second element `&buf[1]` is `(buf + 1)`
 - It can be referenced as `*(buf + 1)` or `buf[1]`


```
int buf[] = {2, 3, 5, 6, 11};
int *p;
p = buf;

*p = *p + 10;
*(p + 1) = *(p + 1) + 10; // {12, 13, 5, 6, 11}
```

1 byte Memory Content
One byte per row



Pointer Arithmetic In Use – C's Performance Focus



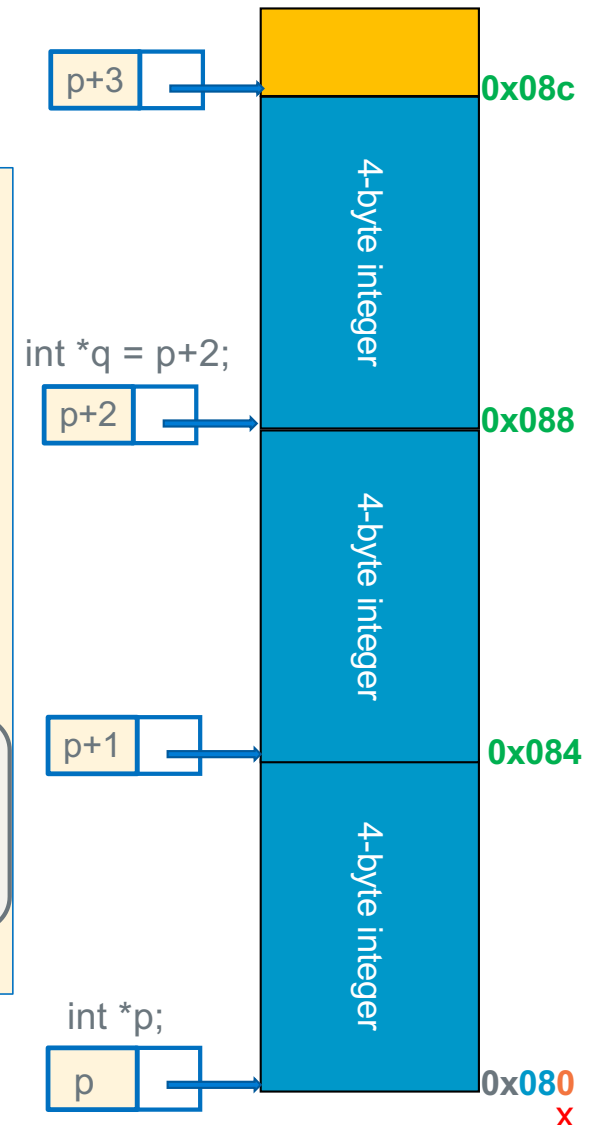
- **Alert!:** C performance focus **does not** perform any array “bounds checking”
- **Performance by Design:** *bound checking **slows down execution** of a properly written program*
- **Example:** array `a` of length `i`, C **does not verify** that `a[j]` or `*(a + j)` is valid (does not check: $0 \leq j < i$)
 - C simply “*translates*” and accesses the memory specified from: `a[j]` to be `*(a + j)` which may be *outside the bounds* of the array
 - OS only “***faults***” for an incorrect access to memory (read-only or not assigned to your process)
 - It does not fault for out of bound indexes or out of scope
- **lack of bound checking** is a **common source of errors and bugs** and is a common criticism of C

Pointer Arithmetic

- You cannot add two pointers (*what is the reason?*)
- A pointer *q* can be subtracted from another pointer *p* when the pointers are the same type – **best done only within arrays!**
- The value of $(p - q)$ is the number of **elements between** the two pointers
 - Using memory address arithmetic (*p* and *q* Rside are both **byte addresses**):

distance in elements = $(p - q) / \text{sizeof}(*p)$

$(p + 3) - p = 3 = (0x08c - 0x080) / 4 = 3$



Pointer Comparisons

- Pointers (**same type**) can be compared with the comparison operators:

<, <=, ==, !=, >=, >

```
int numb[] = {9, 8, 1, 9, 5};
int *end;
int *a;
end = numb + (int) (sizeof(numb)/sizeof(*numb));
a = numb;
while (a < end) // compares two pointers (address)
    /* rest of code including doing an a++ */
```

- Invalid, Undefined, or **risky** pointer arithmetic (some examples)
 - Add, multiply, divide on two pointers
 - Subtract two pointers of different types or pointing at different arrays
 - Compare two pointers of different types
 - Subtract a pointer from an integer

Fast Ways to "Walk" an Array: Use a Limit Pointer

```
int x[] = {0xd4c3b2a1, 0xd4c3b200, 0x12345684};  
int cnt = (int)(sizeof(x) / sizeof(*x));
```

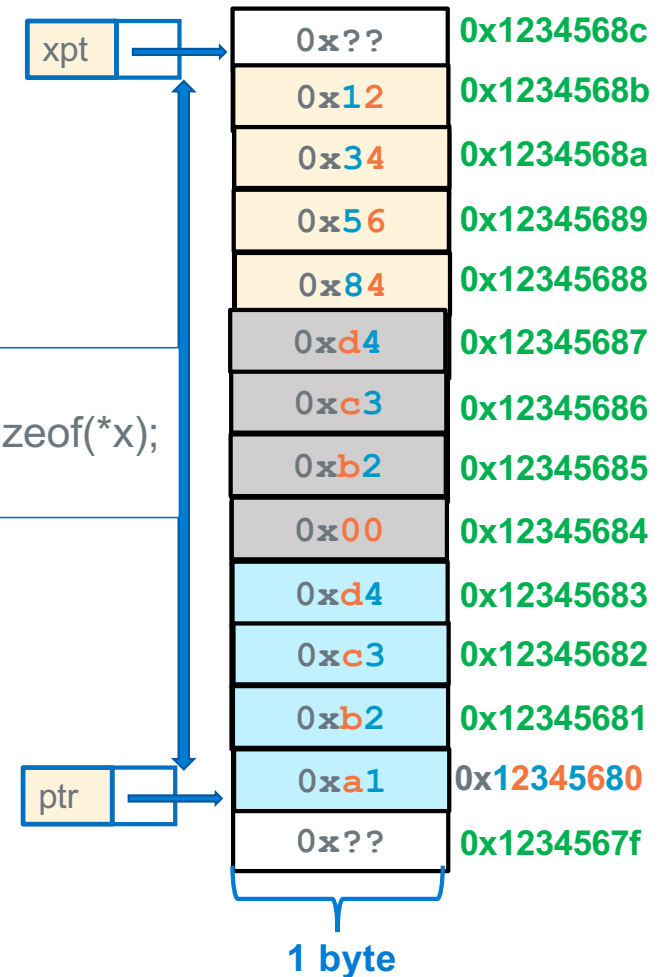
```
int *ptr;  
int *xpt;  
ptr = x; //or &x[0]  
xptr = ptr + cnt;
```

xpt is a loop **limit pointer**
it points 1 element past
the end of the array

```
while (ptr < xpt) {  
    printf("%#x\n", *ptr);  
    ptr++;  
}
```


```
% ./a.out  
0xd4c3b2a1  
0xd4c3b200  
0x12345684
```

```
cnt = 3;  
bytes = cnt * sizeof(*x);  
      = 12
```



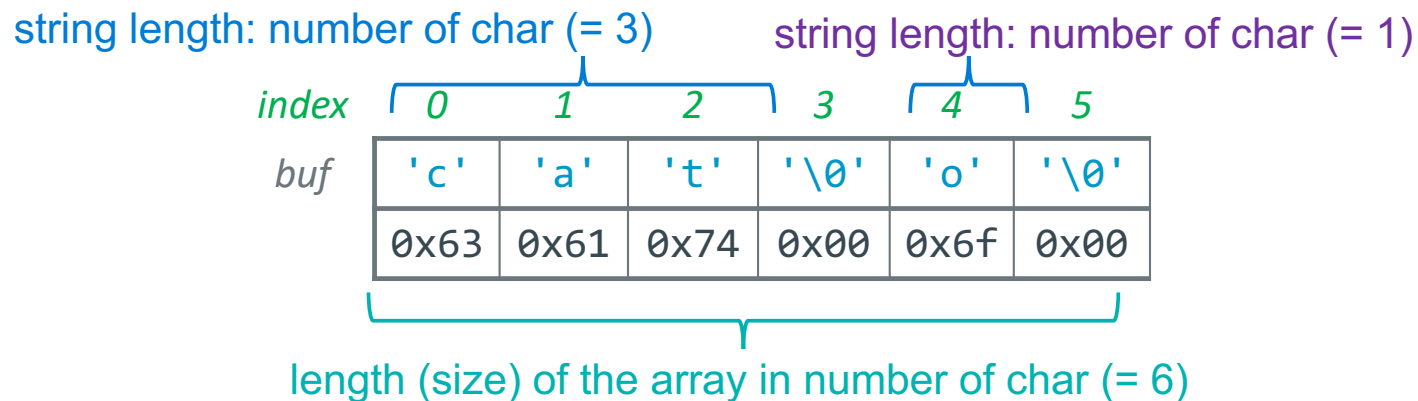
C Strings - 1

- C does not have a **dedicated type** for strings
- Strings are an **array of characters** terminated by a **sentinel termination character**
- `'\0'` is the **Null termination character**; has the **value of zero** (do not confuse with `'0'`)
- An **array of chars** contains **a string only when** it is terminated by a `'\0'`
- **Length of a string** is the number of characters in it, not including the `'\0'`
- Strings in C are **not** objects
 - **No embedded information about them**, you just have a **name** and a **memory location**
 - You **cannot use + or +=** to concatenate strings in C
 - For example, you must **calculate string length** using code at runtime looking for the end



C Strings - 2

- First '`\0`' encountered from the start of the string always indicates the end of a string
- The '`\0`' **does not have to be** in the **last element in the space allocated to the array**
 - But, String length is always less than the size of the array it is contained in
- In the example below, the array `buf` contains two strings
 - One string starts at `&(buf[0])` is `"cat"` with a string length of 3
 - The other string starts at `&(b[4])` is `"o"` with a string length of 1
 - `"o"` has two bytes: `'o'` and `\0`



Defining Strings: Initialization

- When you combine the automatic length definition for arrays with double quote(") **initialization**
 - Compiler automatically adds the null terminator '\0' for you

```
char a[4] = {'c', 'a', 't', '\0'};
char b[] = "cat";
char c[] = {'c', 'a', 't', '\0', 'a', 'b'};
char empty[] = "";
```

// compiler calculates size, adds '\0'
// array size 6, string length 3
// empty string - contains '\0'
// string length = 0

Defining Strings: Initialization Equivalents

- Following definitions create **equivalent** 4-character arrays
 - These are all strings as they all include a null ('\0') terminator

```
char a[4] = {'c', 'a', 't', '\0'};
char b[4] = {'c', 'a', 't', 0};
char c[4] = {'c', 'a', 't'};           // missing initial value defaults to 0
char d[4] = { 99, 97, 116, 0};         // 99 = 'c', 97 = 'a', 116 = 't'
char e[4] = "cat";
char f[4] = "cat\0";                   // literal has 5 chars; array f string
                                       // length is 3
```

Background: Different Ways to Pass Parameters

- **Call-by-reference (or pass by reference)**

- Parameter in the called function is an alias (references the same memory location) for the supplied argument
- Modifying the parameter modifies the calling argument

Call-by-value (or pass by value) (C)

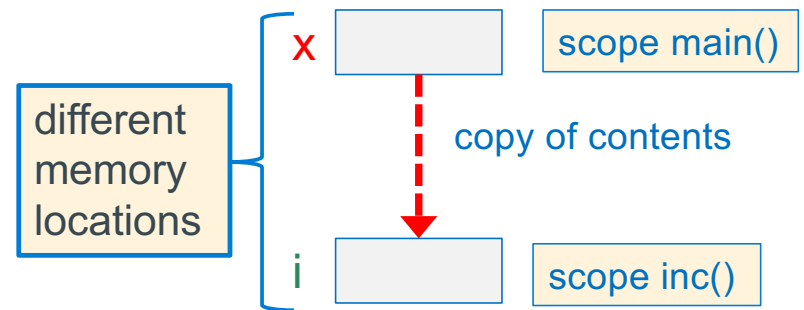
- What **Called** Function Does
 - Passed Parameters are used like local variables
 - Modifying the passed parameter in the function is allowed just like a local variable
 - So, writing to the parameter, only changes the copy
- The return value from a function in C is **by value**

Passing Parameters – Call by Value Example

```
int main(void)
{
    int x = 5;
    inc(x); // makes a copy of x
    printf("%d\n", x); // 5 or 6 ?
}

void inc(int i) // i is local to inc
{
    ++i;
}
```

if this was an expression like `inc(x+1)` it evaluates and stores the result in the memory allocated for the copy



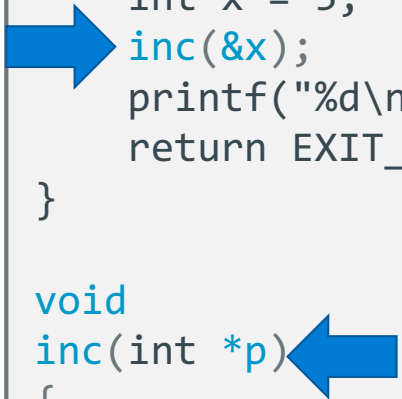
- when `inc(x)` is called, a copy of `x` is made to another memory location
 - `inc()` cannot change the variable `x` since `inc()` does not have the address of `x`, it is local to `main()` so, 5 is printed
- The `inc()` function is free to change its copy of the argument (just like any local variable) remember it does NOT change the parameter in `main()`

Output Parameters

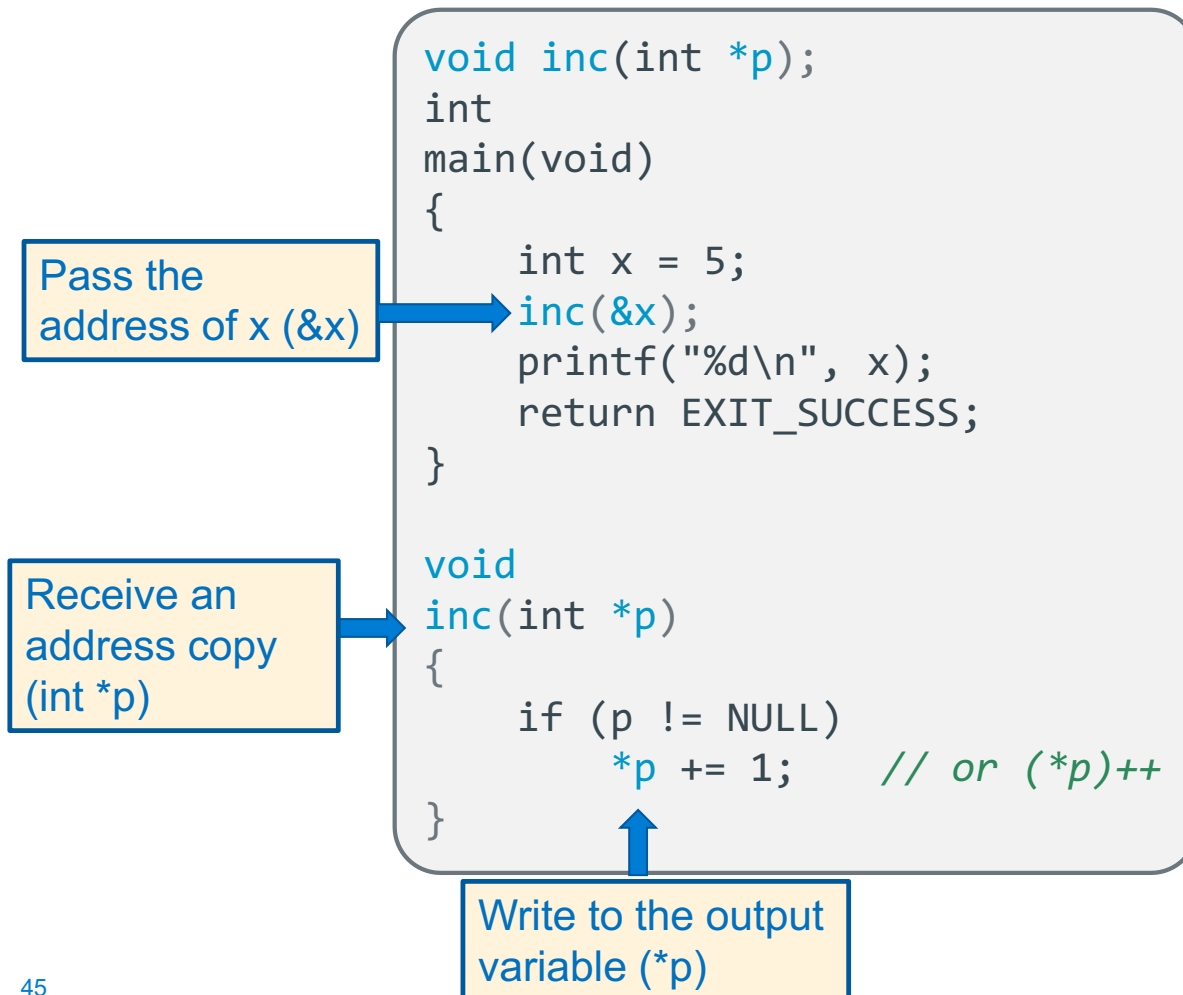
- Passing a **pointer parameter** with the **intent** that the called function will use the address it to store values for use by the **calling function**, then **pointer parameter** is called an **output parameter**
- To pass the address of a variable `x` use the **address operator** (`&x`) or the contents of a pointer variable that points at `x`, or the name of an array (the arrays address)
- To be receive an address in the called function, define the **corresponding parameter type** to be a **pointer**
 - It is common to describe this method as: “**pass a pointer to x**”
- C is still using “*pass by value*”
 - we pass the **value** of the **address/pointer** in a **parameter copy**
 - **The called routine** uses the address to change a variable in the caller's scope

```
void inc(int *p);
int
main(void)
{
    int x = 5;
    inc(&x);
    printf("%d\n", x);
    return EXIT_SUCCESS;
}

void
inc(int *p)
{
    if (p != NULL)
        *p += 1; // or (*p)++
}
```




Example Using Output Parameters



At the Call to inc() in main()

1. Allocate space for p
2. Copy x's address into p



With a pointer to X,

inc() can change x in main()

this is called a side effect

p just like any other local variable

Array Parameters: Call-By-Value or Call-By-Reference?

- `Type[]` array parameter is automatically “promoted” to a pointer of type `Type *`, and a copy of the *pointer* is *passed by value*

```
int main(void)
{
    int numbers[] = {9, 8, 1, 9, 5};

    passa(numbers);
    printf("numbers size:%lu\n", sizeof(numbers)); // 20
    return EXIT_SUCCESS;
}
```

```
void passa(int a[])
{
    printf("a size:%lu\n", sizeof(a)); // 4
    return;
}
```

IMPORTANT:

See the size difference 20 in `main()` in `passa()` is 4 bytes (size of a pointer)

- Call-by-value pointer (callee can change the pointer parameter to point to something else!)
- Acts like call-by-reference (called function can change the contents caller's array)

Arrays As Parameters: What is the size of the array?

- It's tricky to use arrays as parameters, as they are passed as pointers to the start of the array
 - In C, Arrays do not know their own size and at runtime there is no “bounds” checking on indexes

```
int sumAll(int a[]);  
  
int main(void)  
{  
    int numb[] = {9, 8, 1, 9, 5};  
    int sum = sumAll(numb);  
  
    return EXIT_SUCCESS;  
}  
  
int sumAll(int a[])  
{  
    int i, sum = 0;  
    int sz = (int) (sizeof(a)/sizeof(*a));  
    for (i = 0; i < sz; i++) // this does not work  
        sum += a[i];  
}
```

the name is the address, so this is passing a pointer to the start of the array

“inside” the body of sumAll(), the question is: how big is that array? all I have is a POINTER to the first element.....
sz is a 1 on 32 bit arm

Arrays As Parameters, Approach 1: Pass the size

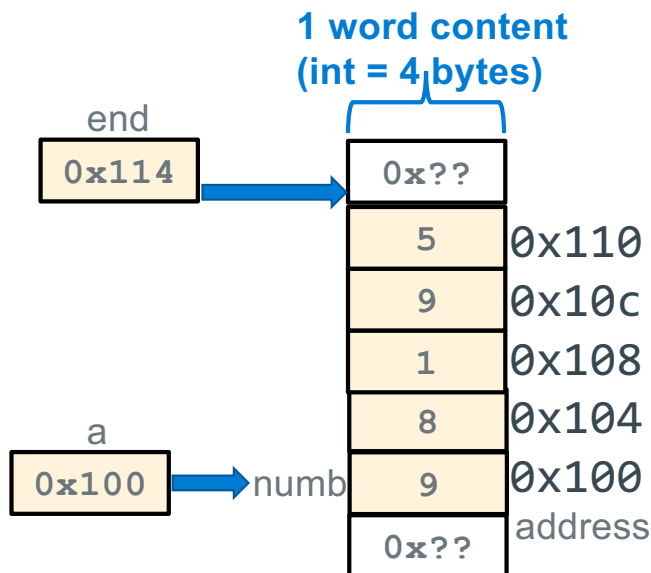
Two ways to pass array size

1. pass the **count** as an additional argument
2. add a **sentinel element** as the last element

remember you can only use `sizeof()` to calculate element count where the array is defined

```
int sumAll(int *a, int size);
int main(void)
{
    int numb[] = {9, 8, 1, 9, 5};
    int cnt = sizeof(numb)/sizeof(numb[0]);

    printf("sum is: %d\n", sumAll(numb, cnt));
    return EXIT_SUCCESS;
}
```



```
int sumAll(int *a, int size)
{
    int sum = 0;
    int *end;
    end = a + size;

    while (a < end)
        sum += *a++;
    return sum;
}
```

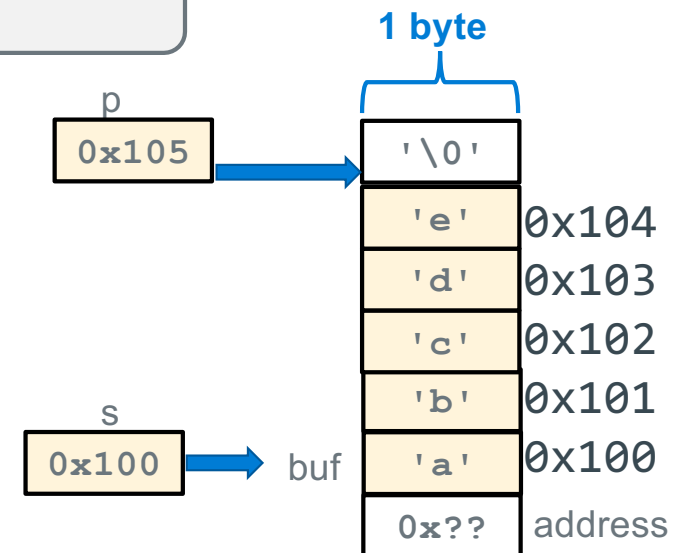
Arrays As Parameters, Approach 2: Use a sentinel element

- A **sentinel** is an element that contains a value that is not part of the normal data range
 - Forms of 0 are often used (like with strings). Examples: '\0', NULL

```
int strlen(char *a);
int main(void)
{
    char buf[] = {'a', 'b', 'c', 'd', 'e', '\0'}; // string

    printf("Number of chars is: %d\n", strlen(buf));
    return EXIT_SUCCESS;
}
```

```
/* Assumes parameter is a terminated string */
int strlen(char *s)
{
    char *p = s;
    if (p == NULL)
        return 0;
    while (*p++)
        ;
    return (p - s - 1);
}
```



Do not overuse strlen()

- C string library function `strlen()` calculates string length **at runtime**
- **Do not overuse `strlen()`, as it walks the array each time called**

```
int count_e(char *s) //  $O(n^2)$  !!!  
{  
    int count = 0;  
    if (s == NULL)  
        return 0;  
    for (int j = 0; j < strlen(s); j++) {  
        if (s[j] == 'e')  
            count++;  
    }  
    return count;  
}
```



```
int count_e(char *s) //  $O(n)$  !!!  
{  
    int count = 0;  
    if (s == NULL)  
        return 0;  
    while (*s) {  
        if (*s++ == 'e')  
            count++;  
    }  
    return count;  
}
```

Comparing strings

- **Characters** can be easily **compared** ($c1 < c2$) as they are numbers, so the **character order** is determined by the ASCII values assigned to each character
 - 65 = A 66 = B 67 = C 68 = D 69 = E 70 = F 71 = G, and so on.
- **Example:** the following strings are in lexicographical (alphabetical) order:
"" "a" "az" "c" "cab" "cabin" "cat" "catastrophe"
- Compare two strings lexicographically (i.e., comparing ASCII values), subtract one from the other

Return Value	Comparison
< 0	$s1 < s2$
> 0	$s1 > s2$
$= 0$	$s1 == s2$

```
int strcmp(char *s1, char *s2)
{
    while (*s1 == *s2) {
        if ((*s1 == '\0') || (*s2 == '\0'))
            break;
        s1++;
        s2++;
    }
    return *s1 - *s2; // character difference
}
```

Copying Strings: Use the Sentinel; libc: strcpy(), strncpy()

- To copy an array, you must copy each character from source to destination array
- Watch overwrites: strcpy assumes the target array size is equal or larger than source array

index	0	1	2	3	4	5
char	'H'	'e'	'l'	'l'	'o'	'\0'

```
char str1[80];  
strcpy(str1, "hello");
```

```
// strncpy adds a length limit on copy  
char str1[6];  
strncpy(str1, "hello", 5); // \0 not copied  
str1[5] = '\0'; // make sure \0 terminated
```

```
char *strcpy(char *s0, char *s1)  
{  
    char *str = s0;  
  
    if ((s0 == NULL) || (s1 == NULL))  
        return NULL;  
    while (*s0++ = *s1++)  
        ;  
    return str;  
}
```

```
char *strncpy(char *s0, char *s1, int len)  
{  
    char *str = s0;  
    if ((s0 == NULL) || (s1 == NULL))  
        return NULL;  
  
    while ((*s0++ = *s1++) && --len)  
        ;  
    return str;  
}
```


Reference: Some String Routines in libc (#include <string.h>)

Function	Description
strlen(<i>str</i>)	returns the # of chars in a C string (before null-terminating character).
strcmp(<i>str1</i> , <i>str2</i>), strncmp(<i>str1</i> , <i>str2</i> , <i>n</i>)	compares two strings; returns 0 if identical, <0 if <i>str1</i> comes before <i>str2</i> in alphabet, >0 if <i>str1</i> comes after <i>str2</i> in alphabet. strncmp stops comparing after at most <i>n</i> characters.
strchr(<i>str</i> , <i>ch</i>) strrchr(<i>str</i> , <i>ch</i>)	character search: returns a pointer to the first occurrence of <i>ch</i> in <i>str</i> , or NULL if <i>ch</i> was not found in <i>str</i> . strrchr find the last occurrence.
strstr(<i>haystack</i> , <i>needle</i>)	string search: returns a pointer to the start of the first occurrence of <i>needle</i> in <i>haystack</i> , or NULL if <i>needle</i> was not found in <i>haystack</i> .
strcpy(<i>dst</i> , <i>src</i>), strncpy(<i>dst</i> , <i>src</i> , <i>n</i>)	copies characters in <i>src</i> to <i>dst</i> , including null-terminating character. Assumes enough space in <i>dst</i> . Strings must not overlap. strncpy stops after at most <i>n</i> chars, and <u>does not</u> add null-terminating char.
strcat(<i>dst</i> , <i>src</i>), strncat(<i>dst</i> , <i>src</i> , <i>n</i>)	concatenate <i>src</i> onto the end of <i>dst</i> . strncat stops concatenating after at most <i>n</i> characters. <u>Always</u> adds a null-terminating character.
strspn(<i>str</i> , <i>accept</i>), strcspn(<i>str</i> , <i>reject</i>)	strspn returns the length of the initial part of <i>str</i> which contains <u>only</u> characters in <i>accept</i> . strcspn returns the length of the initial part of <i>str</i> which does <u>not</u> contain any characters in <i>reject</i> .

2D Array of Char (where elements may contain strings)

- 2D array of chars (where rows may include strings)
- Each row has the same fixed number of memory allocated
- All the rows are the same length regardless of the actual string length)
- The column size must be large enough for the longest string

high memory char aos2d[3][22] = {"my", "two dimensional", "char array"};

aos2d[2]	c	h	a	r		a	r	r	a	y	'\0'											
aos2d[1]	t	w	o		d	i	m	e	n	s	i	o	n	a	l		a	r	r	a	y	'\0'
aos2d[0]	m	y	'\0'																			

low memory

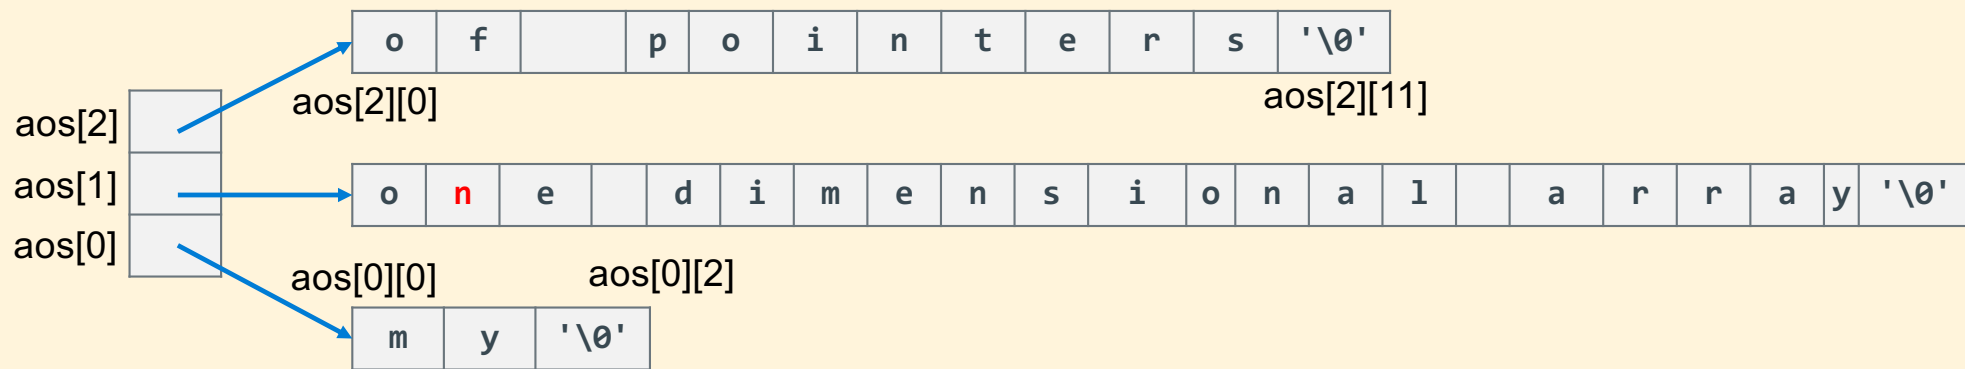
```
#define ROWS 3
char aos[ROWS][22] = { "my", "two dimensional", "char array"};
char (*ptc)[22] = aos; // ptr points at a row of 22 chars

for (int i = 0; i < ROWS; i++)
    printf("%s\n", *(ptc + i));
```

high memory

Pointer Array to Strings (This is NOT a 2D array)

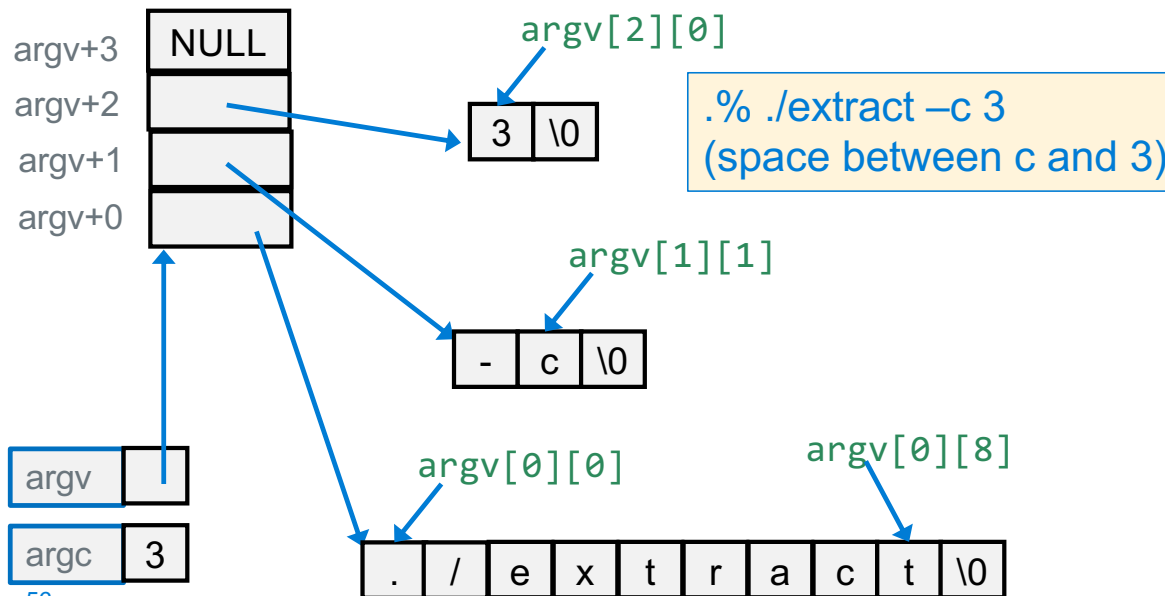
- 2D char arrays are an inefficient way to store strings (wastes memory) unless all the strings are similar lengths, so 2D char arrays are rarely used with string elements
- An array of pointers is common for strings as "rows" can vary in length



- `aos` is an array of pointers; each pointer points at a character array (also a string here)
- Not a 2D array, but any char can be accessed as if it was in a 2D array of chars
 - When I was learning, this was the most confusing syntax aspects of C!


main() Command line arguments: argc, argv

- Arguments are passed to main() as a pointer to an array of pointers (**`**argv`** or **`*argv[]`**)
Conceptually: `% *argv[0] *argv[1] *argv[2] ...`.
- `argc`** is the number of VALID elements (they point at something)
- `*argv`** (**`argv[0]`**) is **usually** the **name** of the executable file (`% ./vim file.c`)
- `*(argv + argc)`** always contains a NULL (0) sentinel
- `*argv[]`** (or **`**argv`**) elements point at **mutable strings!**



```
printf("%s\n", *(argv+0));  
printf("%s\n", *(argv+1));  
printf("%s\n", *(argv+2));
```

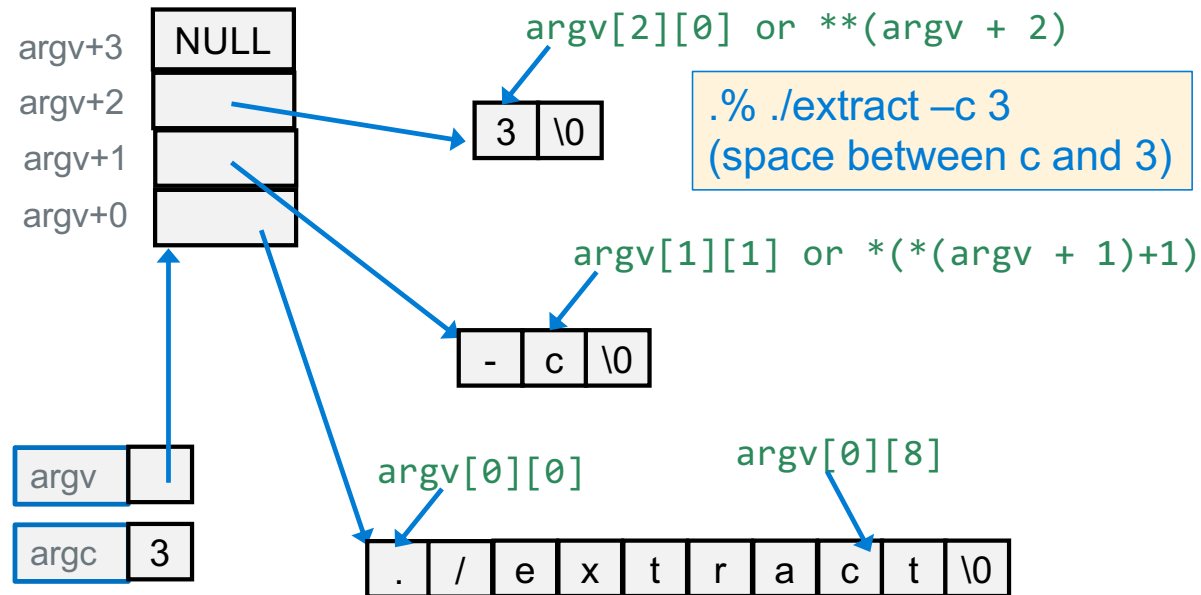
main() Command line arguments: argc, argv



`argv[0][0]` equiv to `** (argv+0)`
`** (argv+0)` equiv `**argv`
`argv[0][8]` equiv `* (*argv + 8)`

```
char *pt = *argv;  
*pt equiv to ** (argv+0)  
*(pt+8) equiv to * (*argv + 8)
```


main() Command line arguments: argc, argv




```
int main(int argc, char *argv[])
{
    for (int i = 0; argv[i] != NULL; i++) {
        for (int j = 0; argv[i][j] != '\0'; j++)
            putchar(argv[i][j]);
        putchar('\n');
    }
    return EXIT_SUCCESS;
}
```

```
int main(int argc, char **argv)
{
    char *pt;
    while ((pt = *argv++) != NULL) {
        while (*pt != '\0')
            putchar(*pt++);
        putchar('\n');
    }
    return EXIT_SUCCESS;
}
```

main() Command line arguments: argc, argv



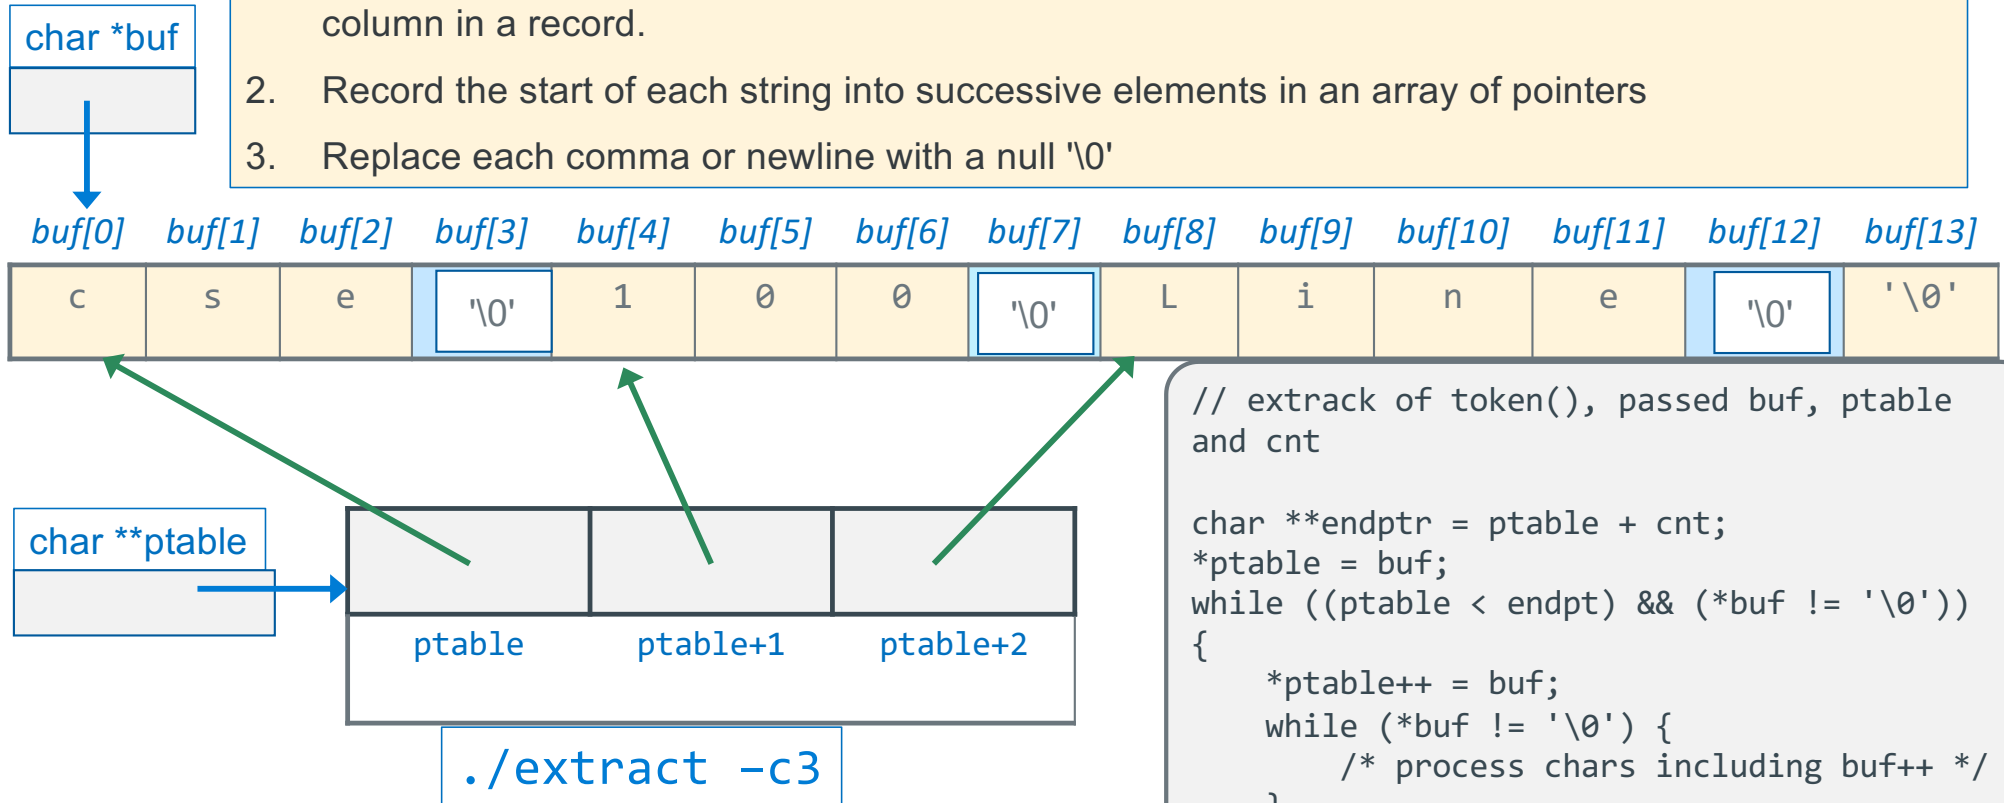
./% ./extract -c 3
(space between c and 3)



./% ./extract -c3
(No space between c and 3)

PA4: Creating a 2D Array of Mutable String Pointers

1. Break a string of comma separated words into individual strings without copying. Do This by walking the string until you see an either a comma , or a newline \n. Each points at a field or column in a record.
2. Record the start of each string into successive elements in an array of pointers
3. Replace each comma or newline with a null '\0'




```
// extract of token(), passed buf, ptable
// and cnt

char **endptr = ptable + cnt;
*ptable = buf;
while ((ptable < endptr) && (*buf != '\0'))
{
    *ptable++ = buf;
    while (*buf != '\0') {
        /* process chars including buf++ */
    }
}
// check for too many or too few fields
```


Review: Pointer Array to Strings

How to access: `aos[1][1]` is `*(*aos + 1) + 1` which contains 'n'
its address is `(*aos + 1) + 1`

aos+2 is not shown due to space limits on the slide



```
char *ptr;  
ptr = *aos;  
*ptr = 'X';  
if (*ptr == ',')  
    or  
if (**aos) == ','
```